

1. INTRODUCTION

1.1 Background

In contemporary healthcare systems, the widespread adoption of Electronic Health Record (EHR) and Electronic Medical Record (EMR) systems has resulted in the generation of massive volumes of clinical documents. Healthcare providers generate extensive unstructured text narratives daily, including admission notes, consultation records, operative reports, discharge summaries, and diagnostic evaluations. Categorizing clinical text into specific medical specialties (e.g., Cardiology, Neurology, Orthopedics, Gastroenterology, Surgery) is a fundamental prerequisite for organizing health databases, routing clinical files, and supporting decision systems.

Traditionally, this categorization relies on manual review by medical coders and health information management staff. As patient volumes increase, manual document indexing becomes increasingly bottlenecked, leading to administrative delays and human classification errors. Natural Language Processing (NLP) and Machine Learning (ML) present viable computational paradigms for automating clinical text categorization, enabling rapid and consistent medical specialty identification.

1.2 Problem Statement

Manual categorization of medical transcription reports into clinical specialties is time-consuming, labor-intensive, and highly prone to inter-annotator variability. Unstructured clinical notes exhibit complex linguistic characteristics, including domain-specific jargon, non-standard medical abbreviations, variable document lengths, and semantic ambiguity. Furthermore, real-world clinical text datasets suffer from severe class imbalance, where dominant medical specialties contain thousands of records, whereas specialized sub-disciplines contain very few. Existing systems lack computationally efficient mechanisms to handle high-dimensional clinical text spaces while remaining lightweight enough for real-time web deployment.

1.3 Need for the System

Automating medical specialty classification yields key operational and clinical benefits:

- *Streamlined Workflow*: Eliminates manual tagging backlogs, allowing clinical reports to

be indexed and routed immediately upon dictation.

- *Enhanced Information Retrieval*: Enables clinicians to query specialty-specific diagnostic records rapidly during emergency care.
- *Reduced Administrative Overhead*: Allows medical records staff to focus on quality assurance rather than manual document sorting.
- *Scalable Health Analytics*: Provides standardized metadata necessary for hospital resource allocation and epidemiological research.
- *Real-Time Accessibility*: Proposes an intuitive web interface for healthcare professionals to upload reports and receive specialty predictions.

1.4 Objectives

The primary objective of this project is to develop and deploy an automated Medical Specialty Classification system using TF-IDF feature representations and an Ensemble Machine Learning architecture. The specific objectives are:

1. To conduct Exploratory Data Analysis (EDA) on the Kaggle Medical Transcriptions (MTSamples) dataset.
2. To build an NLP preprocessing pipeline incorporating lowercasing, punctuation removal, tokenization, stop-word removal, and lemmatization.
3. To extract numerical text features using Term Frequency–Inverse Document Frequency (TF-IDF) vectorization.
4. To train and evaluate candidate machine learning models (Logistic Regression, Support Vector Machine, Random Forest, Multinomial Naive Bayes).
5. To construct a Voting Ensemble Classifier combining predictions from candidate base models.
6. To deploy a real-time, interactive web application using Flask for medical document classification.

1.5 Scope

The scope of this project encompasses:

- *Dataset Scope*: Evaluates 4,999 clinical transcription records across 40 medical specialties from the Kaggle MTSamples dataset.
- *Feature Scope*: Focuses on TF-IDF vectorization derived from the primary narrative column (transcription).
- *Algorithmic Scope*: Evaluates classical ML models (LR, SVM, RF, MNB) and a Voting Ensemble Classifier optimized for CPU deployment.
- *Deployment Scope*: Proposes a lightweight Flask web interface for single-document and batch specialty prediction.

1.6 Project Overview

The project follows a structured engineering workflow:

1. *Data Ingestion & EDA*: Ingesting raw MTSamples CSV data and analyzing feature distributions.
2. *Text Preprocessing*: Removing noise, tokenizing, filtering stop-words, and lemmatizing text.
3. *Feature Engineering*: Transforming clean text into sparse numerical vectors via TF-IDF.
4. *Model Development*: Training baseline classifiers and hyperparameter tuning.
5. *Ensemble Integration*: Combining the outputs of the candidate models using soft voting to generate the final specialty prediction.
6. *Web Deployment*: Serializing model artifacts with Joblib and serving via a proposed Flask web application.

1.7 Expected Outcome

The expected deliverables include:

- A clean, preprocessed clinical text dataset optimized for NLP tasks.
- A proposed Voting Ensemble Classifier targeting robust performance across 40 medical specialties.
- A planned lightweight Flask web application to provide real-time specialty classification.
- Complete technical documentation adhering to MCA departmental standards.

2. SUPPORTING LITERATURE

2.1 Literature Review

A systematic review of contemporary literature was conducted to evaluate existing methodological approaches, feature representations, classification algorithms, and dataset challenges in medical text categorization. The review relies strictly on approved reference literature in clinical natural language processing.

Paper 1 – Omar AM, Elhoseny M, Hassanien AM. Clinical text classification with word representation features and machine learning algorithms. International Journal of Online and Biomedical Engineering (iJOE). 2023.

The paper titled “Clinical Text Classification with Word Representation Features and Machine Learning Algorithms” evaluates the impact of different feature extraction techniques and machine learning algorithms on Electronic Medical Record (EMR) narratives. The authors introduced a four-phase classification framework consisting of text preprocessing, word vectorization, feature reduction, and supervised classification.

To model clinical text narratives, the study compared classical bag-of-words and TF-IDF representations against dense neural embeddings (Word2Vec) across multiple classifiers, establishing baseline benchmarks for clinical text processing. The evaluated classifiers include Logistic Regression, Support Vector Machine (SVM), Naive Bayes, and k-Nearest Neighbors (k-NN).

The empirical results demonstrated that combining Word2Vec embeddings with a k-NN classifier achieved the highest accuracy of 92%. Dense semantic embeddings consistently outperformed sparse unigram Bag-of-Words representations across clinical text categories. However, distance-based k-NN introduces significant computational latency during inference, rendering it inefficient for massive production databases. Furthermore, the study relied on single classifiers without exploring ensemble learning strategies or real-time web deployment frameworks. Crucially, Principal Component Analysis (PCA) was evaluated exclusively in Paper 1 for feature reduction and is not included in our proposed system architecture.

TITLE	Clinical Text Classification with Word Representation Features and Machine Learning Algorithms [1]
AREA OF WORK	Clinical Natural Language Processing / EMR Text Vectorization
DATASET	Electronic Medical Record (EMR) Clinical Transcriptions Dataset
METHODOLOGY / STRATEGY	4-Phase Framework: (1) Preprocessing, (2) Vectorization (BOW, TF-IDF, Word2Vec), (3) Feature Reduction (PCA), (4) Classification
ALGORITHM	Logistic Regression, Support Vector Machine (SVM), Naive Bayes, k-NN, PCA
RESULT / PRECISION	Word2Vec + k-NN achieved the highest classification accuracy of 92%
ADVANTAGES	Proves dense semantic embeddings capture medical context better than unigram BOW
LIMITATIONS	Distance-based k-NN scales poorly; lacks ensemble learning and web UI; PCA used only in literature
RELEVANCE TO MY PROJECT	Establishes the standard 4-phase text preprocessing and feature extraction pipeline
FUTURE PROPOSAL	Develop lightweight ensemble classifiers suitable for real-time web deployment

Table 1: Summary of Paper 1

Paper 2 – Zhang H, Zhu D, Tan H, Shafiq M, Gu Z. Medical specialty classification based on semiadversarial data augmentation. Hindawi Computational Intelligence and Neuroscience. 2023.

The research article titled “Medical Specialty Classification Based on Semiadversarial Data Augmentation” addressed data scarcity and category imbalance in medical specialty classification. The authors developed Semiadversarial Data Augmentation (SemiADA), generating synthetic samples along adversarial attack trajectories to populate sparse decision regions.

Additionally, a Probabilistic Information (PI) layer was introduced post-Softmax to recalculate confidence based on domain noun importance. The paper evaluated a filtered subset of the Kaggle Medical Specialty Classification dataset restricted to 18 medical specialty classes. In contrast, our project utilizes the complete Kaggle Medical Transcriptions (MTSamples) dataset spanning 40 medical specialties.

The proposed SemiADA framework achieved an average performance improvement of 14.9% in accuracy and F1 score over baseline benchmarks. The study proved that domain-specific nouns carry the highest discriminative signal for medical specialty assignment. However, generating synthetic text along adversarial attack paths incurs extreme computational overhead. Fine-tuning heavy transformer models like BERT requires dedicated GPU infrastructure, making real-time inference difficult in resource-constrained web environments.

TITLE	Medical Specialty Classification Based on Semiadversarial Data Augmentation [2]
AREA OF WORK	Clinical Text Data Augmentation / Class Imbalance Handling
DATASET	Kaggle Medical Specialty Classification Dataset (Subset restricted to 18 classes)
METHODOLOGY / STRATEGY	Semiadversarial sample generation along attack paths + Probabilistic Information (PI) noun-weighted Softmax recalculation
ALGORITHM	SemiADA, BERT (Transformer), Softmax with PI Noun Layer
RESULT / PRECISION	Achieved an average 14.9% boost in Accuracy and F1 score over baseline benchmarks
ADVANTAGES	Effectively handles extreme class imbalance using geometry-aware synthetic samples
LIMITATIONS	Extreme computational latency; fine-tuning BERT demands high-end GPU hardware
RELEVANCE TO MY PROJECT	Identifies class imbalance in MTSamples data; motivates TF-IDF + Ensemble ML alternative
FUTURE PROPOSAL	Combine TF-IDF term weighting with ensemble learning to classify 40 specialties on standard CPUs

Table 2: Summary of Paper 2

Paper 3 – Blanchard AE, Gao S, Yoon HJ, Tourassi G. A keyword-enhanced approach to handle class imbalance in clinical text classification. IEEE Journal of Biomedical and Health Informatics. 2022.

Blanchard et al., writing in the *IEEE Journal of Biomedical and Health Informatics*, focused on preventing deep neural models from overfitting on rare clinical categories in long pathology narratives. The authors introduced a data-centric strategy by injecting short class-associated

keyword sequences into mini-batches during neural network training, formulating a combined loss function

$$L = L_{\text{docs}} + \alpha L_{\text{key}}$$

The study utilized the National Cancer Institute (NCI) Surveillance, Epidemiology, and End Results (SEER) cancer pathology registry dataset. The approach evaluated Convolutional Neural Networks (CNNs), keyword-constrained loss functions, and standard data resampling baselines (oversampling and undersampling).

Injecting keyword constraints significantly improved macro F1 score and recall on rare classes. The study demonstrated that overall micro-accuracy is a deceptive metric in clinical NLP because high overall scores can mask near-zero recall on rare clinical conditions. However, the approach requires manual curation of comprehensive keyword dictionaries by clinical domain experts for every target category, limiting automated scalability to new medical domains.

TITLE	A Keyword-Enhanced Approach to Handle Class Imbalance in Clinical Text Classification [3]
AREA OF WORK	Imbalanced Clinical Narrative Classification / Cancer Registry NLP
DATASET	NCI SEER Surveillance Cancer Pathology Registry Reports Dataset
METHODOLOGY / STRATEGY	Ingesting keyword sequences into mini-batches during training to constrain the neural loss function
ALGORITHM	Convolutional Neural Networks (CNNs), Keyword-Constrained Loss Optimization
RESULT / PRECISION	Significantly improved macro F1 score and recall on rare classes without degrading majority performance
ADVANTAGES	Prevents deep models from memorizing noisy tokens in long text documents with rare labels
LIMITATIONS	Requires manual domain expert curation of keyword dictionaries for every target category
RELEVANCE TO MY PROJECT	Demonstrates that keyword/n-gram features preserve rare class signals, supporting TF-IDF
FUTURE PROPOSAL	Utilize automated TF-IDF n-gram feature extraction without requiring manual dictionary creation

Table 3: Summary of Paper 3

2.2 Literature Review Summary

Table 4 presents a concise summary of the reviewed literature alongside our proposed system.

Paper	Area	Dataset	Algorithm	Advantages	Relation to Proposed Work
Omar et al. [1]	Clinical Vectorization	EMR Text	LR, SVM, NB, k-NN	High accuracy (92%)	Establishes 4-phase preprocessing pipeline
Zhang et al. [2]	Data Augmentation	MTSamples (18 classes)	SemiADA, BERT	+14.9% accuracy gain	Highlights MTSamples class imbalance
Blanchard et al. [3]	Class Imbalance	NCI SEER Pathology	CNN + Keyword Loss	High macro F1 on rare classes	Justifies TF-IDF term weighting
Proposed Work	Specialty Categorization	MTSamples (40 classes)	TF-IDF + Voting Ensemble	Fast, lightweight, web UI	Complete MCA Project System

Table 4: Overall Summary of all paper

Evolution of Research

Clinical text processing research has transitioned from basic Bag-of-Words word counting toward neural word embeddings (Word2Vec), transformer architectures (BERT), and complex adversarial data augmentation (SemiADA). While deep neural networks achieve strong performance, their high computational cost and memory footprints prevent real-time deployment in standard healthcare web environments.

Identified Research Gaps

The literature review reveals three primary research gaps:

1. **High Computational Overhead:** Existing deep learning systems require dedicated GPUs and introduce substantial inference latency.
2. **Lack of Real-Time Web Deployment:** Prior studies focus exclusively on offline evalua-

tion metrics without building accessible user interfaces.

3. **Single Model Reliance:** Most prior works rely on single classifiers rather than leveraging multi-model ensemble diversity.

Motivation for Proposed System

Our project addresses these research gaps by combining TF-IDF vectorization with a Voting Ensemble Classifier (aggregating Logistic Regression, Linear SVM, Random Forest, and Multinomial Naive Bayes) intended for deployment via a lightweight Flask web application. This architecture is designed to target accurate classification across all 40 medical specialties in the full MTSamples dataset while aiming for efficient CPU inference for real-world clinical use.

2.3 Findings and Proposals

2.3.1 Findings

The review of the existing literature shows that machine learning techniques can effectively classify medical and clinical text into different medical specialties. Traditional machine learning algorithms such as Logistic Regression, Support Vector Machine, Random Forest, and Naive Bayes have demonstrated good performance in text classification tasks when appropriate feature extraction techniques are used.

The literature also indicates that TF-IDF feature extraction is effective for converting unstructured clinical text into numerical feature vectors. It helps identify important medical terms and provides suitable high-dimensional input for machine learning algorithms.

Another important finding is that the performance of individual machine learning models may vary depending on the characteristics of the dataset. Some algorithms perform better for certain classes, while others provide better overall generalization. Therefore, relying on a single classifier may not always produce the best prediction performance.

The reviewed studies further show that ensemble learning techniques, which combine predictions from multiple machine learning models, can improve the robustness and reliability of classification systems. However, challenges such as class imbalance, noisy clinical text, missing values, and variations in medical terminology need to be considered during data preprocessing. Based on these findings, it is observed that combining effective text preprocessing, TF-IDF fea-

ture extraction, multiple machine learning algorithms, and an ensemble classification approach can provide a suitable solution for medical specialty classification.

2.3.2 Proposal

Based on the findings from the literature review, the proposed system aims to develop a Medical Specialty Classification system using TF-IDF and Ensemble Machine Learning.

The system will use clinical transcription data containing medical reports and their corresponding medical specialties. Initially, the dataset will be analysed and preprocessed by handling missing values, removing unnecessary information, converting text into a consistent format, and cleaning the clinical transcription data.

After preprocessing, the clinical text will be transformed into numerical feature vectors using the TF-IDF Vectorizer. Multiple machine learning algorithms, including Logistic Regression, Linear Support Vector Machine, Random Forest, and Multinomial Naive Bayes, will then be trained using the extracted features.

The predictions generated by the individual machine learning models will be combined using an ensemble learning approach, such as Voting Classification. The proposed approach is expected to improve the reliability and classification performance compared with depending on a single machine learning model.

Finally, the trained system will accept a clinical transcription as input and predict the most appropriate medical specialty. The system will be evaluated using suitable performance measures such as accuracy, precision, recall, and F1-score.

Thus, the proposed system combines text preprocessing, TF-IDF feature extraction, multiple machine learning algorithms, and ensemble learning to provide an efficient approach for the automatic classification of medical clinical transcriptions.

3. SYSTEM ANALYSIS

System analysis establishes the functional requirements, dataset architecture, preprocessing steps, exploratory data analysis, and algorithmic framework for the medical specialty classification system.

The Kaggle Medical Transcriptions (MTSamples) dataset, sourced from Kaggle [4], provides clinical narrative dictations across multiple healthcare sub-disciplines, enabling comprehensive evaluation of natural language processing and machine learning models for medical document categorization.

Dataset Link: <https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions>

Size: The dataset comprises 4,999 clinical records spanning 40 distinct medical specialties.

Data Structure: It contains six attributes including row index (Unnamed: 0), short report abstract (description), target medical department (medical_specialty), sample dictation title (sample_name), main clinical narrative (transcription), and extracted key terms (keywords).

Sparsity: The primary text column (transcription) is highly complete, containing 33 missing records (0.66%), which are dropped during preprocessing. The target label (medical_specialty) contains zero missing values.

3.1 EXPLORATORY DATA ANALYSIS

3.1.1 ABOUT THE DATASET

Dataset Overview

The Kaggle Medical Transcriptions dataset provides authentic clinical narratives dictating patient histories, diagnostic evaluations, physical examinations, and surgical procedures:

- Total Records: 4,999 clinical reports.
- Total Features: 5 independent variables + 1 target variable (medical_specialty).
- Data Types: Textual narratives (transcription, description, keywords), categorical labels (medical_specialty), and integer index.
- Missing Values: Present in description (6 records), transcription (33 records), and keywords (1,068 records).

Table 5 summarizes the core dataset attributes.

Property	Description
Dataset Name	Kaggle Medical Transcriptions (MTSamples) Dataset
Source	Kaggle (https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions)
Total Records	4,999 clinical reports
Number of Columns	6 fields (Unnamed: 0, description, medical_specialty, sample_name, transcription, keywords)
Input Variable	transcription (Primary clinical narrative text)
Target Variable	medical_specialty (40 unique clinical specialties)
Advantages	Authentic clinical dictations; complete target label annotations (0 missing labels); broad multi-class domain coverage
Limitations	Severe class imbalance across categories; 33 missing transcriptions; 1,068 missing keyword entries; variable text length

Table 5: Characteristics of the Kaggle Medical Transcriptions Dataset

The dataset attributes are categorized as follows:

1. **Primary Clinical Narrative** (transcription): The central predictor attribute containing unstructured, dictation-style medical reports that detail symptoms, findings, diagnoses, and treatments.
2. **Clinical Abstract** (description): Short one-to-two sentence summary of the clinical consultation or procedure.
3. **Extracted Keywords** (keywords): Comma-delimited clinical terms and anatomical references extracted from the narrative.
4. **Sample Identification** (sample_name): Dictation title indicating the type of procedure or medical note.
5. **Target Variable** (medical_specialty): The multi-class categorical label representing the specialized department (40 clinical specialties).

3.1.2 STATISTICAL ANALYSIS

Data exploration was conducted using Python data analysis tools (pandas, numpy). Inspection via shape, info(), dtypes, and describe() confirmed that the dataset contains 4,999 records and 6 columns without structural corruption. Table 6 describes each column, its data type, and missing value counts.

Column Name	Data Type	Description	Missing Count
Unnamed: 0	Integer	Record Identifier / Row Index	0
description	Text	Short summary of clinical report	6
medical_specialty	Categorical	Target specialty label (40 classes)	0
sample_name	Text	Report title / dictation sample title	0
transcription	Text	Primary clinical narrative (Input)	33
keywords	Text	Extracted key terms	1,068

Table 6: Dataset Column Descriptions and Missing Value Audit

Descriptive statistical analysis revealed significant variation in class distribution across the 40 medical specialties. Table 7 lists the record counts for the top 10 majority specialties alongside the 10 rarest minority specialties.

Top 10 Specialties	Count	Bottom 10 Specialties	Count
Surgery	1,103	Cosmetic / Plastic Surgery	27
Consult - History and Phy.	516	Pediatric / Neonatal	22
Cardiovascular / Pulmonary	372	Sleep Medicine	20
Orthopedic	355	Bariatrics	18
Radiology	273	Speech - Language	9
General Medicine	259	Lab Medicine - Pathology	8
Gastroenterology	230	Allergy / Immunology	7
Neurology	223	Hospice - Palliative Care	6
SOAP / Chart / Progress Notes	166	Executive Evaluation	2
Obstetrics / Gynecology	160	Autopsy	2

Table 7: Distribution of Top 10 Majority and Bottom 10 Minority Medical Specialties

Statistical analysis of narrative text lengths indicated a positively skewed distribution. While the median document length is approximately 350 to 500 words, certain comprehensive opera-

tive notes and consultative evaluations extend beyond 2,000 words. This variation in document length further supports the use of TF-IDF feature representation with vector normalization for consistent numerical feature scaling.

3.2 DATA VISUALIZATION

Data visualization is an important stage of Exploratory Data Analysis (EDA) that helps in understanding the characteristics, distributions, and relationships among clinical narratives and medical specialties. Visual representations provide valuable insights into the dataset, reveal hidden patterns, and assist in verifying the suitability of the text data for machine learning classification.

Missing Value Analysis

Figure 1 illustrates the distribution of missing values across dataset attributes. As shown, the primary target attribute (`medical_specialty`) contains zero missing values. The primary input text narrative (`transcription`) contains 33 missing records, while the auxiliary keywords attribute exhibits 1,068 missing records.




Figure 1: Missing Values by Dataset Attribute

Top 10 Medical Specialties

Figure 2 displays the record counts for the top 10 medical specialties in the dataset. Surgery represents the single largest category with 1,103 clinical reports, followed by Consult - History and Phy. (516 reports), Cardiovascular / Pulmonary (372 reports), and Orthopedic (355 reports).




Figure 2: Top 10 Medical Specialties by Number of Transcriptions

Complete Specialty Distribution

Figure 3 illustrates the full class distribution across all 40 medical specialties in the MTSamples corpus. The chart demonstrates severe class imbalance, ranging from majority classes such as Surgery (1,103 records) down to highly specialized minority classes such as Hospice - Palliative Care, Allergy / Immunology, Lab Medicine - Pathology, and Autopsy with under 10 records each.




Figure 3: Distribution of Medical Specialties

Transcription Length Distribution

Figure 4 displays the word count distribution for clinical narratives, fitted with a Kernel Density Estimate (KDE) curve. The majority of medical transcriptions range between 200 and 600

words, exhibiting a right-skewed distribution with extended narratives reaching up to 3,000 words.




Figure 4: Distribution of Clinical Transcription Length

Transcription Length Across Top Specialties

Figure 5 presents box plots comparing narrative word count distributions across the top 10 medical specialties. Specialties such as Radiology and Gastroenterology feature shorter median narrative lengths (250 words), whereas Orthopedic, Surgery, and Consult - History and Phy. contain longer document narratives (median 500 words) with multiple high-length outliers exceeding 1,500 to 2,500 words.




Figure 5: Transcription Length Distribution Across Top 10 Medical Specialties

Medical Word Cloud Visualization

Figure 6 displays a word cloud of the most frequent clinical terms extracted from the medical transcription corpus. High-frequency domain terms include "placed", "well", "used", "time",

3.3 DATA PREPROCESSING

Data preprocessing is an essential stage in developing a machine learning model, as it transforms raw, unstructured clinical narratives into a structured and meaningful format suitable for mathematical modeling. In this project, the original dataset consisted of free-text clinical dictations containing clinical terminology, punctuation, formatting artifacts, and varying narrative lengths. The preprocessing stage included data cleaning, text normalization, analysis of feature variables, and examination of the target class distribution.

3.3.1 DATA CLEANING

Text preprocessing is essential in clinical NLP to remove noisy elements, normalize medical vocabulary, and transform raw narrative text into clean input suitable for mathematical vectorization. The cleaning pipeline executes eight sequential operations:

1. *Removing Missing Values*: Records containing empty or null transcription fields (33 records) are dropped from the dataset.
2. *Removing Duplicates*: Verification ensures no duplicate clinical reports exist in the corpus.
3. *Lowercasing*: All narrative text is converted to lowercase to standardize character representations.
4. *Removing Punctuation*: Commas, periods, semicolons, and special punctuation symbols are stripped.
5. *Removing Special Characters*: Numeric digits and non-alphanumeric symbols are removed.
6. *Tokenization*: Free-text narratives are segmented into individual word tokens.
7. *Stop-Word Removal*: Generic English stop-words are filtered out while preserving clinical vocabulary.
8. *Lemmatization*: Word tokens are mapped to their canonical base form using the WordNet Lemmatizer.

Table 8 provides concrete text transformation examples for each step.

Preprocessing Step	Input Text Example	Processed Output Example
Raw Input Text	"SUBJECTIVE: Patient suffers from acute chest pain!"	"SUBJECTIVE: Patient suffers from acute chest pain!"
1. Missing Values	Checked: Non-null text	Retained
2. Duplicates	Checked: Unique entry	Retained
3. Lowercasing	"SUBJECTIVE: Patient suffers..."	"subjective: patient suffers from acute chest pain!"
4. Punctuation	"subjective: patient suffers..."	"subjective patient suffers from acute chest pain"
5. Special Chars	"chest pain! 100%"	"chest pain"
6. Tokenization	"patient suffers pain"	['patient', 'suffers', 'pain']
7. Stop-Words	['patient', 'suffers', 'from', 'pain']	['patient', 'suffers', 'pain']
8. Lemmatization	['patient', 'suffers', 'pain']	['patient', 'suffer', 'pain']

Table 8: Step-by-Step Clinical Text Preprocessing Examples

3.3.2 ANALYSIS OF FEATURE VARIABLES

The transcription field is selected as the primary feature variable because it contains the complete clinical narrative. Other metadata fields, such as description and keywords, are not included in the primary classification feature set because the project focuses on full transcription-based specialty classification and the keywords field contains substantial missing values. Unstructured clinical narratives are converted into high-dimensional numerical feature vectors using Term Frequency–Inverse Document Frequency (TF-IDF) vectorization, which extracts unigram and bigram tokens while applying sublinear frequency scaling and L_2 normalization.

3.3.3 ANALYSIS OF CLASS VARIABLE

The target variable `medical_specialty` contains 40 distinct medical specialty categories. The class distribution exhibits severe imbalance, ranging from 1,103 reports in Surgery to only 2 reports in Autopsy and Executive Evaluation. To ensure fair multi-class evaluation, stratified train-test splitting (80:20) is utilized to preserve class proportions, and ensemble voting consensus is deployed to avoid model bias toward majority classes.

3.4 ANALYSIS OF ALGORITHMS

This project employs supervised machine learning algorithms to classify clinical transcription narratives into 40 distinct medical specialties. Multiple classification algorithms are considered to evaluate their effectiveness in distinguishing domain-specific clinical vocabulary based on extracted TF-IDF feature vectors. The proposed algorithms include Support Vector Machine (Linear SVM), Random Forest (RF), Logistic Regression (LR), and Multinomial Naive Bayes (MNB), integrated through a Voting Ensemble Classifier. These algorithms were chosen because they offer strong predictive performance in high-dimensional sparse text domains, provide architectural diversity, and execute efficiently on standard computing hardware without requiring specialized GPU accelerators.

3.4.1 DETAILED STUDY REPORT OF THE PROPOSED ALGORITHMS

The proposed system evaluates diverse machine learning estimators alongside statistical feature extraction to construct a robust specialty classification engine. Each algorithm offers distinct theoretical properties, mathematical formulations, and operational characteristics when applied to high-dimensional clinical text spaces.

1. Support Vector Machine (SVM)

Support Vector Machine (SVM) is a robust supervised learning algorithm founded on statistical learning theory and structural risk minimization principles. It constructs an optimal decision boundary (hyperplane) in a high-dimensional feature space to separate target classes while maximizing the geometric margin between them. Linear SVM (Linear SVC) is exceptionally well suited for text categorization where sparse TF-IDF vectors offer natural linear separability.

A. Theoretical Foundation and Optimization Model Given a training dataset of clinical transcription records:

$$D = \{(x_1, y_1), (x_2, y_2), \dots, (x_N, y_N)\} \quad (1)$$

where $x_i \in \mathbb{R}^d$ represents the sparse TF-IDF feature vector across d vocabulary terms and $y_i \in \{1, 2, \dots, K\}$ denotes the target medical specialty ($K = 40$), the linear decision hyperplane is defined as:

$$w^T x + b = 0 \quad (2)$$

where $w \in \mathbb{R}^d$ is the weight vector normal to the hyperplane and $b \in \mathbb{R}$ is the bias scalar. To handle real-world clinical vocabulary noise and lexical overlap between specialties, a soft-margin optimization formulation is applied using non-negative slack variables $\xi_i \geq 0$:

$$\min_{w,b,\xi} \frac{1}{2} \|w\|^2 + C \sum_{i=1}^N \xi_i \quad (3)$$

subject to the linear constraints:

$$y_{i,k}(w_k^T x_i + b_k) \geq 1 - \xi_i, \quad \xi_i \geq 0, \quad \forall i = 1, \dots, N \quad (4)$$

where $C > 0$ is the regularization parameter controlling the trade-off between maximizing the geometric margin width ($\frac{2}{\|w\|}$) and minimizing training classification error.

B. Dual Formulation and Multi-Class Decision Function By introducing Lagrange multipliers $\alpha_i \geq 0$, the primal optimization problem is transformed into its Wolfe dual representation:

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j K(x_i, x_j) \quad (5)$$

subject to $0 \leq \alpha_i \leq C$ and $\sum_{i=1}^N \alpha_i y_i = 0$. For high-dimensional sparse text vectors, the Linear Kernel $K(x_i, x_j) = x_i^T x_j$ is utilized.

To support 40 medical specialties, a One-vs-Rest (OvR) strategy is implemented, training 40 independent binary classifiers. The decision function evaluating an incoming clinical transcription x is:

$$f_k(x) = w_k^T x + b_k \quad (6)$$

The predicted medical specialty $\hat{y}$ corresponds to the class with the maximum decision score:

$$\hat{y} = \arg \max_{k \in \{1, \dots, K\}} (w_k^T x + b_k) \quad (7)$$

C. Hyperparameter Tuning and Clinical Relevance The predictive performance of Linear SVM depends primarily on optimizing the regularization parameter C . Linear SVM excels in clinical transcription classification due to its immunity to high-dimensional overfitting, low memory footprint during inference, and reliance solely on critical support vectors along deci-

sion margins.

D. Probability Calibration for Soft-Voting Integration Standard Linear Support Vector Classifiers (LinearSVC) compute uncalibrated signed geometric distances to separating hyperplanes, which do not directly provide posterior probability distributions (`predict_proba`). Because the proposed ensemble relies on soft-voting probability aggregation, the Linear SVM decision scores are calibrated to produce class probability estimates before soft-voting aggregation. This is achieved using Scikit-learn’s `CalibratedClassifierCV`, which applies Platt scaling via an internal cross-validated logistic sigmoid transformation:

$$P(y = k | x) = \frac{1}{1 + \exp(A \cdot f_k(x) + B)} \quad (8)$$

where scalar parameters A and B are estimated through maximum likelihood on internal validation folds. Probability calibration maps decision boundary margins monotonically into valid multi-class probability distributions, enabling Linear SVM to participate effectively in soft-voting probability aggregation alongside inherently probabilistic estimators.




Figure 7: Activity Diagram of Support Vector Machine

2. Random Forest (RF)

Random Forest is a high-performance ensemble learning algorithm based on decision tree aggregation. It constructs a multitude of uncorrelated decision trees during training and outputs the ensemble consensus. By combining Bootstrap Aggregation (Bagging) with random feature subspace selection, Random Forest mitigates variance without increasing prediction bias.

A. Architectural Mechanics and Ensemble Formulation Given a training dataset of N clinical transcription records, Random Forest generates B bootstrap samples by sampling with replacement. For each bootstrap sample, an individual decision tree $T_b(x)$ is trained. At each split node of every tree, a random subset of $m = \sqrt{d}$ TF-IDF features is evaluated to identify the optimal split.

The final specialty class prediction for an input vector x is derived through majority voting:

$$\hat{y} = \text{mode}(T_1(x), T_2(x), \dots, T_B(x)) \quad (9)$$

Alternatively, posterior class probabilities are computed as the average over all B decision trees:

$$P(y = k | x) = \frac{1}{B} \sum_{b=1}^B P_b(y = k | x) \quad (10)$$

B. Node Splitting Criteria and Impurity Measures Decision tree node splits are evaluated using impurity metrics such as Gini Impurity or Information Gain (Entropy).

- **Gini Impurity (G):**

$$G = 1 - \sum_{k=1}^K p_k^2 \quad (11)$$

where p_k represents the proportion of instances belonging to specialty class $k \in \{1, \dots, 40\}$ at a given node.

- **Entropy (H) and Information Gain (IG):**

$$H(S) = - \sum_{k=1}^K p_k \log_2(p_k) \quad (12)$$

$$IG(S, A) = H(S) - \sum_{v \in \text{Values}(A)} \frac{|S_v|}{|S|} H(S_v) \quad (13)$$

C. Feature Importance Analysis and Clinical Interpretability A major benefit of Random Forest is its capacity to quantify feature importance using Mean Decrease in Impurity (MDI). The importance index $I(A_j)$ of TF-IDF feature A_j across all B decision trees is computed as:

$$I(A_j) = \frac{1}{B} \sum_{b=1}^B \sum_{n \in N_b(A_j)} \Delta G(n) \quad (14)$$

where $\Delta G(n)$ is the Gini impurity reduction achieved at node n by splitting on feature A_j . In clinical NLP, this enables identification and ranking of the most discriminative medical keywords for each specialty.

D. Hyperparameter Optimization The primary hyperparameters tuned for Random Forest include:

- `n_estimators` (B): Number of decision trees (typically 100 to 300).
- `max_depth`: Maximum depth of decision trees to prevent single-tree memorization.
- `min_samples_split`: Minimum samples required to split an internal node.
- `min_samples_leaf`: Minimum samples required at a terminal leaf node.

Random Forest captures complex non-linear keyword interactions that linear models may overlook, adding structural diversity to the ensemble.




Figure 8: Activity Diagram of Random Forest

3. Logistic Regression (LR)

Logistic Regression is a foundational statistical learning algorithm for classification problems. Belonging to the family of Generalized Linear Models (GLMs), it is extended to multi-class

classification using the multinomial Softmax formulation, modeling posterior probabilities directly.

A. Probabilistic Model and Multinomial Softmax Transformation For an input TF-IDF vector $x \in \mathbb{R}^d$ and $K = 40$ medical specialty classes, the algorithm computes linear logit scores $z_k = w_k^T x + b_k$. The posterior probability that transcription x belongs to specialty k is given by the Softmax function:

$$P(y = k | x) = \frac{\exp(w_k^T x + b_k)}{\sum_{j=1}^K \exp(w_j^T x + b_j)} \quad (15)$$

where $w_k \in \mathbb{R}^d$ is the weight vector associated with specialty k and $b_k \in \mathbb{R}$ is the class-specific bias term.

B. Maximum Likelihood Estimation and Loss Function Model parameter estimation is performed by minimizing the Categorical Cross-Entropy (Negative Log-Likelihood) loss function over N training samples, augmented with an L_2 Ridge penalty:

$$J(W, b) = -\frac{1}{N} \sum_{i=1}^N \sum_{k=1}^K \mathbb{I}(y_i = k) \ln P(y_i = k | x_i) + \frac{\lambda}{2} \sum_{k=1}^K \|w_k\|_2^2 \quad (16)$$

where $\mathbb{I}(\cdot)$ is the indicator function, and $\lambda > 0$ controls the regularisation strength to prevent coefficient explosion on sparse text inputs. Parameters are optimized via gradient descent algorithms:

$$w_k^{(t+1)} = w_k^{(t)} - \alpha \nabla_{w_k} J(W, b) \quad (17)$$

C. Clinical Relevance and Interpretability Logistic Regression serves as an essential baseline model due to its high interpretability. The exponentiated weight coefficients $\exp(w_{k,j})$ represent odds ratios for clinical terms, providing quantitative insight into which medical keywords escalate the probability of assigning a transcription to a specific specialty.




Figure 9: Activity Diagram of Logistic Regression

4. Multinomial Naive Bayes (MNB)

Multinomial Naive Bayes is a specialized probabilistic classification algorithm rooted in Bayes' Theorem, specifically formulated for discrete word counts and term frequency distributions in document categorization tasks.

A. Probabilistic Foundation and Bayes' Theorem The algorithm computes the posterior probability of clinical transcription $x = (x_1, \dots, x_m)$ belonging to medical specialty class c using Bayes' rule alongside the assumption of conditional feature independence given the class label:

$$P(c | x) \propto P(c) \prod_{i=1}^m P(x_i | c) \quad (18)$$

where $P(c)$ is the prior probability of specialty c (derived from training class proportions) and $P(x_i | c)$ is the class-conditional feature likelihood.

B. Feature Likelihood and Laplace Smoothing To prevent zero-probability penalties for clinical terms absent in the training split of a specific specialty, Laplace (additive) smoothing is applied:

$$P(x_i | c) = \frac{N_{ci} + \alpha}{N_c + \alpha|V|} \quad (19)$$

where N_{ci} is the cumulative occurrence count of term x_i in specialty class c , N_c is the total token count across all records in class c , $|V|$ is the vocabulary size, and $\alpha = 1.0$ is the smoothing hyperparameter.

C. Computational Efficiency and Role in Medical Text Classification Multinomial Naive Bayes executes in linear time with minimal memory overhead. It functions reliably on smaller clinical training subsets, providing an efficient baseline that contributes probabilistic diversity to the ensemble.




Figure 10: Activity Diagram of Multinomial Naive Bayes

5. TF-IDF Feature Vectorizer

Term Frequency–Inverse Document Frequency (TF-IDF) is a statistical vectorization technique that converts unstructured clinical text narratives into high-dimensional numerical feature vectors. In our proposed medical specialty classification system, TF-IDF operates as the core feature engineering component.

A. Statistical Foundation and Mathematical Formulation The vectorization process evaluates the relative importance of each term t within a specific document d against the entire clinical corpus D . Term Frequency, denoted as $\text{TF}(t, d)$, measures word occurrences scaled sublinearly:

$$\text{TF}(t, d) = 1 + \log(f_{t,d}) \quad (\text{if } f_{t,d} > 0) \quad (20)$$

Inverse Document Frequency, denoted as $\text{IDF}(t, D)$, computes the logarithmic inverse ratio:

$$\text{IDF}(t, D) = \log \left(\frac{1 + |D|}{1 + |\{d \in D : t \in d\}|} \right) + 1 \quad (21)$$

The composite TF-IDF weight is:

$$\text{TF-IDF}(t, d, D) = \text{TF}(t, d) \times \text{IDF}(t, D) \quad (22)$$

To ensure comparability across documents of varying lengths, Euclidean L_2 normalization is applied:

$$v_{\text{norm}} = \frac{v}{\|v\|_2} = \frac{v}{\sqrt{\sum_i v_i^2}} \quad (23)$$

B. N-Gram Extraction and Sparsity Optimization The vectorizer extracts unigrams and bigrams, capturing both isolated medical terms (e.g., "appendectomy") and multi-word phrases (e.g., "coronary artery"). A minimum document frequency threshold filters out rare typos, yielding a normalized, highly efficient sparse matrix.

C. Clinical Term Weighting and Down-Weighting Common Words Through this formulation, terms appearing across almost all reports receive low weights, whereas specialty-specific diagnostic vocabulary receives high discriminative weights.




Figure 11: Activity Diagram of TF-IDF Feature Vectorizer

6. Voting Ensemble Classifier

The Voting Ensemble Classifier is the final meta-classification mechanism in the proposed system. Rather than relying on a single learning algorithm that may exhibit inductive biases across imbalanced specialties, the ensemble combines predictions from diverse candidate models (Logistic Regression, Linear SVM, Random Forest, and Multinomial Naive Bayes) to generate a consensus decision.

A. Ensemble Mechanics and Soft-Voting Aggregation Soft voting requires calibrated class probability estimates ($P(y = k | x)$) across all candidate base estimators. While Logistic Regression, Random Forest, and Multinomial Naive Bayes intrinsically output class probabilities, standard Linear SVM outputs uncalibrated margin distances. Therefore, the Linear SVM decision scores are calibrated using `CalibratedClassifierCV` to produce reliable posterior probability estimates before soft-voting aggregation. In the soft-voting configuration, each calibrated base classifier $m \in \{1, \dots, M\}$ generates a posterior class probability distribution $P_m(y = k | x)$ indicating the likelihood that a given clinical transcription x belongs to medical specialty k . The ensemble then computes the arithmetic mean of these posterior probabilities across all M estimators:

$$P_{\text{ensemble}}(y = k | x) = \frac{1}{M} \sum_{m=1}^M P_m(y = k | x) \quad (24)$$

B. Final Multi-Class Specialty Prediction The final predicted medical specialty $\hat{y}$ is selected as the category that maximizes the aggregated ensemble probability:

$$\hat{y} = \arg \max_{k \in \{1, \dots, K\}} P_{\text{ensemble}}(y = k | x) \quad (25)$$

C. Variance Reduction and Synergy of Diverse Estimators By integrating linear probabilistic modeling (Logistic Regression), maximum-margin geometric boundaries (Linear SVM), non-linear decision trees (Random Forest), and conditional frequency modeling (Multinomial Naive Bayes), the ensemble offsets the individual weaknesses of each classifier, reduces predictive variance, and improves overall classification stability across all 40 medical specialties.




Figure 12: Activity Diagram of Voting Ensemble Classifier

Working of the Proposed Algorithms

The proposed machine learning classification framework will execute through the following nine sequential steps:

1. **Dataset Ingestion:** Ingesting the raw Kaggle MTSamples dataset containing 4,999 clinical records across 40 medical specialties.
2. **Data Cleaning & Null Audit:** Auditing and removing records with missing clinical

transcriptions (33 null records), handling missing keywords (1,068 records), and standardizing text casing.

3. **Stratified Train-Test Splitting:** Partitioning the cleaned dataset into an 80% training set and a 20% testing set using stratified sampling prior to feature extraction to strictly prevent data leakage and preserve class distributions across all 40 medical specialties.
4. **Text Preprocessing Pipeline:** Applying the modular 5-step NLP preprocessing sequence: character cleaning, lowercasing, tokenization, domain stop-word filtering, and WordNet morphological lemmatization.
5. **TF-IDF Feature Representation:** Fitting the TF-IDF vectorizer strictly on the training partition with sublinear scaling and unigrams/bigrams, then transforming both training and test partitions into high-dimensional, L_2 -normalized sparse feature matrices.
6. **Independent Training of Candidate Classifiers:** Fitting candidate supervised learning algorithms (Support Vector Machine with probability calibration, Random Forest, Logistic Regression, and Multinomial Naive Bayes) on the training features.
7. **Hyperparameter Optimization:** Tuning regularisation parameters, smoothing factors, and estimators via grid search cross-validation.
8. **Voting Ensemble Aggregation:** Combining calibrated posterior class probability distributions generated by the candidate classifiers through soft-voting consensus aggregation via average pooling.
9. **Medical Specialty Prediction & Deployment:** Evaluating test clinical transcriptions across multi-class metrics, serializing the pipeline with Joblib, and integrating with a lightweight Flask web application for real-time specialty predictions and confidence scores.

3.5 PROJECT PIPELINE

The proposed Medical Specialty Classification system follows a structured pipeline that transforms raw clinical transcription narratives into reliable medical specialty predictions using machine learning. The pipeline begins with clinical text data acquisition, followed by data cleaning, stratified train-test splitting, text preprocessing, TF-IDF feature extraction, machine

learning model development and voting ensemble aggregation, and finally medical specialty prediction through a user-friendly Flask web application. Each stage plays an important role in ensuring accurate and reliable classification across all 40 medical specialties without data leakage.




Figure 13: Project Pipeline Diagram of Medical Specialty Classification System

The project pipeline consists of the following stages:

1. **Dataset:** The Kaggle MTSamples Medical Transcription Dataset containing 4,966 cleaned clinical reports across 40 medical specialties is used as the input corpus.
2. **Exploratory Data Analysis (EDA):** The distribution of transcription lengths, medical specialty frequencies, vocabulary counts, and class imbalances is systematically analyzed and statistically profiled.
3. **Data Cleaning & Pruning:** Raw clinical notes are audited for null values; 33 records with missing transcription narratives are removed, leaving 4,966 high-quality reports.

4. **Stratified Train–Test Split:** The cleaned dataset is partitioned using stratified sampling into an 80% training partition and a 20% testing partition prior to vectorization to strictly prevent test-set data leakage.
5. **Text Preprocessing:** Training and testing clinical narratives undergo character cleaning, lowercasing, tokenization, medical stop-word filtering, and WordNet morphological lemmatization.
6. **TF-IDF Feature Extraction:** A TF-IDF vectorizer is fitted exclusively on the preprocessed training partition and subsequently transforms both training and testing narratives into normalized sparse numerical feature matrices.
7. **Machine Learning Model Development:** Multiple supervised machine learning algorithms, including Calibrated Support Vector Machine (SVM), Random Forest (RF), Logistic Regression (LR), and Multinomial Naive Bayes (MNB), are trained on the training feature matrix.
8. **Hyperparameter Tuning:** Grid search cross-validation is performed to optimize critical hyperparameters across candidate estimators.
9. **Soft Voting Ensemble Formulation:** A Soft Voting Ensemble is formulated using average probability aggregation to synthesize calibrated class probability distributions from all candidate models, mitigating individual estimator variance.
10. **Model Evaluation and Comparison:** All trained candidate models and the Soft Voting Ensemble are rigorously evaluated using standard multi-class classification metrics: Accuracy, Precision, Recall, and F1-Scores.
11. **Model Selection and Serialization:** The finalized pipeline—combining the fitted TF-IDF vectorizer and calibrated Soft Voting Ensemble—is serialized using Joblib for deployment.
12. **Deployment and Prediction:** The serialized model pipeline is integrated into a lightweight Flask web application, enabling healthcare professionals to enter or upload patient clinical notes and receive real-time specialty predictions accompanied by classification confidence scores.

3.6 Project Timeline

The project timeline outlines the development schedule, review presentations, and sprint releases for the MCA Mini Project.

Period / Date	Planned Activity
17.07.2026	Project Proposal & Synopsis Approval by Guide
20.07.2026 – 21.07.2026	Project Proposal Presentation before Approval Committee
Week 1–2	Dataset Collection, EDA, and Data Cleaning
Week 3–4	Text Preprocessing and WordNet Lemmatization
Week 5	TF-IDF Feature Representation & Vectorization
08.09.2026	First Project Presentation (Phase 1 Faculty Review)
09.09.2026	Sprint Release I (Dataset Exploration & Architecture)
Week 6	Candidate ML Model Training (Calibrated SVM, RF, LR, MNB)
18.09.2026	Sprint Release II
Week 7	Hyperparameter Tuning and Candidate Model Comparison
Week 8	Soft Voting Ensemble Formulation & Evaluation
29.09.2026 – 30.09.2026	Interim Project Presentation
Week 9	Flask Web Integration & Serialized Model Deployment
09.10.2026	Sprint Release III
Week 10–11	Multi-Class Evaluation, Confusion Matrix, and Testing
22.10.2026 – 23.10.2026	Final Project Presentation
30.10.2026	Final Report Submission

Table 9: Project Timeline

3.7 FEASIBILITY ANALYSIS

Feasibility analysis is conducted to evaluate the practicality of developing and implementing the proposed Medical Specialty Classification system. It assesses whether the project can be successfully developed using the available computational resources, technologies, budget, and time. The feasibility study ensures that the proposed solution is technically viable, economi-

cally affordable, operationally practical, and achievable within the designated project schedule.

3.7.1 Technical Feasibility

The proposed system is technically feasible because it utilizes widely adopted open-source software, established Natural Language Processing (NLP) libraries, and standard machine learning toolkits. The development environment consists of Python programming language, Google Colab, and Visual Studio Code, while libraries such as NLTK, Scikit-learn, Pandas, NumPy, Matplotlib, and Seaborn provide comprehensive functionality for text preprocessing, feature extraction, model training, and performance evaluation. The final system will be deployed using the Flask WSGI web framework, enabling healthcare practitioners to enter or upload clinical notes and obtain specialty predictions through an intuitive web interface. The required software tools are freely available and can be executed on standard CPU-based computing systems without requiring specialized high-performance GPU hardware.

3.7.2 Economic Feasibility

The proposed system is economically feasible because it relies entirely on open-source software and publicly available resources. The Kaggle MTSamples clinical transcription dataset used in this project is freely accessible, and all development tools, including Python, Google Colab, Flask, and the required machine learning libraries, are available without commercial licensing costs. Furthermore, because the classical machine learning workflow (TF-IDF paired with Linear SVM, Logistic Regression, Random Forest, and Naive Bayes) is computationally efficient, it eliminates the need for expensive cloud GPU instances or specialized hardware accelerators. Consequently, the overall development expenditure is minimal, making the project exceptionally cost-effective and suitable for academic research and clinical prototype deployment.

3.7.3 Operational Feasibility

The proposed system is operationally feasible because it is designed to provide a simple and user-friendly interface for medical specialty prediction. Healthcare professionals and administrative personnel will be able to input clinical transcription narratives directly through the web application, after which the system will automatically preprocess the text, extract TF-IDF feature vectors, and generate specialty predictions using the trained Voting Ensemble model. The

straightforward workflow minimizes user effort while providing an efficient, reliable, and accessible method for preliminary clinical categorization and automated medical workflow triaging.

3.7.4 Schedule Feasibility

The project is considered schedule feasible because its development activities are planned in a structured and phased manner. The initial phases, including dataset acquisition, exploratory data analysis, data cleaning, and text preprocessing, have been planned and initiated according to the project schedule. The subsequent phases, such as machine learning model development, probability calibration, voting ensemble aggregation, performance evaluation, Flask web application deployment, and system testing, are scheduled sequentially according to the project timeline. The availability of suitable development tools, pre-existing datasets, and established algorithms supports the timely completion of the project within the academic schedule.

3.8 SYSTEM ENVIRONMENT

The system environment specifies the software tools, libraries, and hardware resources required for the development and implementation of the proposed Medical Specialty Classification system. The project uses standard software tools and computing resources for text processing, feature extraction, machine learning model development, performance evaluation, and web application deployment.

3.8.1 Software Environment

The software environment consists of Python and commonly used open-source libraries for Natural Language Processing, data analysis, machine learning, visualization, and web application development. Google Colab and Visual Studio Code are used for development, experimentation, and modular scripting, while Flask is used for deploying the trained model pipeline as an interactive web application. These tools are freely available and do not require specialized commercial software licenses. Table 10 details each software tool, its type, and its specific purpose in the proposed system.

Software / Tool	Type	Purpose
Python	Programming Language	Primary programming language used for data preprocessing, feature extraction, model training, and backend scripts.
Google Colab	Cloud IDE	Cloud-based interactive environment used for exploratory data analysis, model experimentation, and visualization.
Visual Studio Code	Development IDE	Primary IDE for modular Python script development, project organization, and debugging.
L ^A T _E X	Typesetting System	Document preparation system used for compiling academic reports and formatted tables.
Pandas	Data Manipulation	Handling structured CSV data, dataset filtering, null audits, and category mapping.
NumPy	Numerical Computing	Efficient multi-dimensional array manipulation and mathematical operations.
NLTK	NLP Toolkit	Text tokenization, stop-word elimination, and WordNet morphological lemmatization.
Scikit-learn	Machine Learning	TF-IDF feature extraction, individual classifier training, hyperparameter tuning, and Voting Ensemble aggregation.
Matplotlib	Data Visualization	Rendering statistical charts, class distribution plots, and performance curves.
Seaborn	Statistical Graphics	High-level statistical visualization for categorical frequencies and confusion matrices.
WordCloud	Text Visualization	Generating lexical frequency word clouds from clinical narratives.
Flask	Web Framework	Lightweight WSGI framework for serving serialized model predictions via an interactive web interface.

Table 10: Software Environment and Tool Specifications

3.8.2 Hardware Environment

The proposed Medical Specialty Classification system does not require specialized hardware for its development or operation. The classical machine learning approach, combining TF-IDF feature extraction with ensemble supervised classifiers, is computationally lightweight compared to deep learning transformers. A standard laptop or desktop computer with sufficient processing capability, memory, and storage can be used for text preprocessing, feature extraction, model training, and Flask web application hosting. Table 11 outlines the minimum and recommended hardware configurations required for developing and running the proposed system.

Hardware Component	Minimum Requirement	Recommended Specification
Processor (CPU)	Dual-Core Processor (2.0 GHz or higher)	Quad-Core / Octa-Core Intel Core i5/i7 or AMD Ryzen 5/7 (Multi-core CPU; no GPU required)
RAM	8 GB DDR4	16 GB DDR4/DDR5 (Supports parallel cross-validation and multiple candidate estimators)
Storage	20 GB available disk space	256 GB / 512 GB Solid State Drive (SSD) for rapid dataset read/write and virtual environments
Operating System	Windows 10 (64-bit)	Windows 10/11 (64-bit), Ubuntu Linux 20.04+, or macOS
Internet Connection	Standard Broadband (5 Mbps)	High-Speed Broadband Connection (For package installation, dataset download, and cloud backup)

Table 11: Hardware Specifications and Environment Requirements

4. SYSTEM DESIGN

System design translates the analytical requirements and clinical preprocessing pipelines into a robust, deployable machine learning architecture. This section presents the formal design, modular implementation, algorithmic mechanics, and data-partitioning protocols designed to classify unstructured medical transcription narratives into clinical specialties.

4.1 MODEL(S) BUILDING

The classification framework employs a multi-model paradigm comprising four distinct supervised learning families: linear margin optimization (Support Vector Machine), probabilistic link modeling (Logistic Regression), non-parametric bagging (Random Forest), and Bayesian conditional independence (Multinomial Naive Bayes). Furthermore, an Ensemble Voting meta-classifier is synthesized to consolidate individual algorithmic strengths into an optimal consensus predictor.

4.1.1 Implementation Code

The model training and inference pipeline is implemented in Python using `scikit-learn` and `imbalanced-learn`. The system executes sublinear TF-IDF vectorization, cost-sensitive loss calibration, probability calibration, and ensemble voting aggregation. The core implementation code of the model building pipeline is presented below:

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression
from sklearn.svm import LinearSVC
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import accuracy_score, classification_report

# 1. Load Filtered 8-Specialty Dataset
df = pd.read_csv('filtered_mtsamples_8classes.csv')
X = df['clinical_text'].astype(str)
```

```
y = df['medical_specialty']

# 2. Stratified 80/20 Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.20, random_state=42, stratify=y
)

# 3. Controlled TF-IDF Feature Extraction (8,000 Dimensions)
tfidf = TfidfVectorizer(
    max_features=8000, ngram_range=(1, 2),
    sublinear_tf=True, min_df=2, max_df=0.7, norm='l2'
)
X_train_tfidf = tfidf.fit_transform(X_train)
X_test_tfidf = tfidf.transform(X_test)

# 4. Base Classifiers with Class Balancing
lr = LogisticRegression(class_weight='balanced', max_iter=1000, random_state=42)
svm = LinearSVC(class_weight='balanced', max_iter=3000, random_state=42)
cal_svm = CalibratedClassifierCV(LinearSVC(class_weight='balanced', max_iter=3000))
rf = RandomForestClassifier(n_estimators=100, max_depth=25,
                           class_weight='balanced', random_state=42)
mnb = MultinomialNB(alpha=0.5)

# 5. Soft Voting Ensemble (Weighted Consensus)
voting_ensemble = VotingClassifier(
    estimators=[('lr', lr), ('svm', cal_svm), ('rf', rf), ('mnb', mnb)],
    voting='soft', weights=[2, 2, 1, 1], n_jobs=1
)

# 6. Model Training & Evaluation
voting_ensemble.fit(X_train_tfidf, y_train)
y_pred = voting_ensemble.predict(X_test_tfidf)

print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```

4.1.2 Model Planning (Training/Testing Split of the Dataset)

Model planning governs the experimental protocol to guarantee unbiased evaluation and eliminate data leakage. The dataset consists of 1,663 curated clinical transcriptions spanning 8 mutually exclusive clinical departments.

To preserve the underlying specialty distribution while evaluating generalizability on unseen patient records, an **80:20 Stratified Partitioning Protocol** was established:

- **Training Partition (80%)**: Comprises 1,330 clinical records used exclusively for vocabulary extraction, inverse document frequency parameter estimation, and classifier weight optimization.
- **Testing Partition (20%)**: Comprises 333 clinical records completely held out during feature extraction and model fitting, reserved strictly for final inference validation.
- **Stratification Protocol**: The split is governed by the class distribution vector y , ensuring that minority specialties (such as Ophthalmology and ENT) maintain proportional representation across both partitions.
- **Data Leakage Prevention**: The TF-IDF vectorizer parameters ($IDF(t)$, vocabulary dictionary, document frequencies) are fit *strictly* on X_{train} . The test partition X_{test} is transformed using the fixed vocabulary parameters learned from training.

Table 12 details the exact sample distribution across the training and testing sets for each of the 8 clinical specialties.

Clinical Specialty	Total Records	Train Samples (80%)	Test Samples (20%)	Class Proportion
Cardiovascular / Pulmonary	371	297	74	22.31%
Orthopedic	355	284	71	21.35%
Gastroenterology	224	179	45	13.47%
Neurology	223	178	45	13.41%
Urology	156	125	31	9.38%
Obstetrics / Gynecology	155	124	31	9.32%
ENT - Otolaryngology	96	77	19	5.77%
Ophthalmology	83	66	17	4.99%
Total Corpus	1,663	1,330	333	100.00%

Table 12: Stratified Training and Testing Split Distribution across 8 Clinical Specialties

4.1.3 Training

Build the Model Using the Train Dataset

During the training phase, the four candidate classifiers and the meta-voting ensemble are parameterized using the 1,330 training instances vectorized across the 8,000-dimensional TF-IDF space.

1. **Logistic Regression (LR) Optimization:** Trained as a multinomial log-linear classifier minimizing the penalized cross-entropy objective:

$$\mathcal{L}_{\text{LR}}(\mathbf{W}) = - \sum_{i=1}^N \sum_{k=1}^K w_k \cdot y_{i,k} \log P(y_i = k \mid \mathbf{x}_i; \mathbf{W}) + \frac{1}{2C} \|\mathbf{W}\|_F^2 \quad (26)$$

where $w_k = \frac{N}{K \cdot N_k}$ represents the class-balancing weight inversely proportional to class frequency N_k , penalizing classification errors on under-represented specialties such as Ophthalmology and ENT. Optimization is resolved via the L-BFGS quasi-Newton solver over 1,000 iterations.

2. **Support Vector Machine (Linear SVM) Optimization:** Trained as an unconstrained maximum-margin formulation using one-vs-rest (OvR) decomposition. The primal optimization objective is formulated as:

$$\begin{aligned} \min_{\mathbf{w}_k, b_k, \xi} \quad & \frac{1}{2} \|\mathbf{w}_k\|_2^2 + C \sum_{i=1}^N w_{y_i} \xi_{i,k} \\ \text{subject to} \quad & y_{i,k} (\mathbf{w}_k^T \mathbf{x}_i + b_k) \geq 1 - \xi_{i,k}, \quad \xi_{i,k} \geq 0 \end{aligned} \quad (27)$$

The linear kernel is computationally optimal for sparse, high-dimensional textual matrices ($D = 8,000$), resisting empirical overfitting. To allow probability integration within the Soft Voting Ensemble, a duplicate SVM model is calibrated via Platt scaling using 5-fold cross-validation (CalibratedClassifierCV).

3. **Random Forest (RF) Optimization:** Constructed as an ensemble of $B = 100$ bootstrap decision trees constrained to maximum depth $d_{\text{max}} = 25$. At each split node, a random subset of $m = \sqrt{8,000} \approx 89$ features is evaluated against the Gini impurity criterion:

$$I_G(p) = 1 - \sum_{k=1}^K p_k^2 \quad (28)$$

Balanced bootstrap subsampling (`class_weight='balanced'`) dynamically recalculates sample weights for each bootstrap tree, preventing dominant surgical and cardiovascular notes from monopolizing the split nodes.

4. **Multinomial Naive Bayes (MNB) Optimization:** Computes prior probabilities $P(C_k)$ and feature likelihood parameters $\theta_{k,j} = P(w_j | C_k)$ from training count matrices with Laplace smoothing parameter $\alpha = 0.5$:

$$\theta_{k,j} = \frac{N_{k,j} + \alpha}{N_k + \alpha \cdot |V|} \quad (29)$$

where $N_{k,j}$ is the aggregate TF-IDF weight of word j within specialty k , and $|V| = 8,000$ represents the vocabulary size.

5. **Voting Ensemble Consensus Synthesis:** The predictions of the four trained estimators are pooled via two ensemble voting schemes:

- **Hard Voting:** Assigns the predicted specialty by majority vote:

$$\hat{y}_{\text{hard}} = \operatorname{argmax}_k \sum_{m=1}^M w_m \cdot \mathbb{I}(\hat{y}_m = k) \quad (30)$$

- **Soft Voting:** Averages the posterior probability distributions across models weighted by empirical confidence:

$$\hat{y}_{\text{soft}} = \operatorname{argmax}_k \sum_{m=1}^M w_m \cdot P_m(y = k | \mathbf{x}) \quad (31)$$

where weights are assigned as $w_{\text{LR}} = 2$, $w_{\text{SVM}} = 2$, $w_{\text{RF}} = 1$, and $w_{\text{MNB}} = 1$.

4.1.4 Testing

Evaluate the Performance of the Model Using the Test Dataset

Testing executes a strictly separated inference pipeline on the 333 unobserved clinical narratives in the testing partition X_{test} :

1. **Inference Pipeline Flow:** Each test record undergoes identical preprocessing transformations: lowercasing, punctuation stripping, stop-word elimination, and projection into

the 8,000-dimensional vector space using the pre-fitted vectorizer:

$$\mathbf{x}_{\text{test}} = \text{TFIDF}(t_{\text{test}}; \mathbf{V}_{\text{train}}) \quad (32)$$

2. **Independent Prediction Generation:** The vector is fed into each standalone model and the voting ensemble pipelines:

$$\hat{\mathbf{y}}_{\text{test}} = \mathcal{M}(\mathbf{X}_{\text{test_tfidf}}) \quad (33)$$

3. **Metric Evaluation Protocol:** Predicted labels $\hat{\mathbf{y}}_{\text{test}}$ are benchmarked against true ground-truth labels $\mathbf{y}_{\text{test}}$. Performance is quantified across global exact accuracy, macro-averaged precision, macro-averaged recall, class-specific F1-scores, and weighted F1-metrics to rigorously verify robustness against remaining class skewness.

5. RESULTS AND DISCUSSION

This section delivers an exhaustive empirical evaluation of the developed medical specialty classification models on the held-out test partition of the MTSamples dataset. It presents comparative performance metrics, mathematical evaluation formulations, confusion matrix diagnostics, class-level performance statistics, prediction analysis on unseen clinical records, and an analytical discussion of the findings.

5.1 COMPARISON OF MODELS

The four supervised base estimators (Logistic Regression, Support Vector Machine, Random Forest, and Multinomial Naive Bayes) alongside the two Voting Ensemble configurations (Hard Voting and Soft Voting) were evaluated on the 333 unobserved clinical test records.

Table 13 presents the quantitative comparison across five core evaluation dimensions: Classification Accuracy, Macro-Averaged Precision, Macro-Averaged Recall, Macro-Averaged F1-Score, Weighted F1-Score, and Training Execution Time.

Model Architecture	Accuracy	Macro Prec	Macro Recall	Macro F1	Weighted F1	Train Time (s)
Support Vector Machine (SVM)	0.8859	0.9051	0.8902	0.8971	0.8862	0.12
Soft Voting Ensemble (4 Models)	0.8799	0.9041	0.8795	0.8898	0.8818	2.49
Logistic Regression (LR)	0.8709	0.8966	0.8768	0.8846	0.8730	0.37
Hard Voting Ensemble (4 Models)	0.8709	0.8966	0.8768	0.8846	0.8730	2.01
Random Forest (RF)	0.8198	0.8340	0.8257	0.8274	0.8209	1.54
Multinomial Naive Bayes (MNB)	0.7748	0.8602	0.7331	0.7743	0.7824	0.01

Table 13: Empirical Performance Comparison of Base Models and Voting Ensembles (8 Specialties)

As illustrated in Table 13, **every single candidate model exceeds the minimum 75.00% performance threshold:**

- **Support Vector Machine (SVM)** achieved the highest individual classification accuracy of **88.59%** with a Macro F1-score of **0.8971**, demonstrating the exceptional capacity of linear maximum-margin hyperplanes to separate high-dimensional textual feature matrices.
- **Soft Voting Ensemble** delivered the most resilient generalized performance, achieving **87.99% accuracy**, **90.41% macro precision**, and **0.8898 macro F1-score**, successfully

mitigating individual model variances by pooling calibrated probability vectors.

- **Logistic Regression (LR)** followed closely with **87.09% accuracy** and **0.8846 Macro F1**, training in just 0.37 seconds.
- **Random Forest (RF)** achieved **81.98% accuracy**, completely reversing the initial base-line collapse (7.85%) through constrained tree depth and controlled 8,000-dimensional TF-IDF features.
- **Multinomial Naive Bayes (MNB)** achieved **77.48% accuracy** and 86.02% precision with near-instantaneous execution time (0.01 seconds).

Figure 14 visually contrasts the accuracy, macro F1, and macro recall metrics across all models against the red benchmark target line of 75%.




Figure 14: Model Performance Comparison across Accuracy, Macro F1, and Macro Recall against the 75% Target Threshold

5.2 EVALUATION METRICS

To quantify multi-class text categorization under class imbalance, five standard statistical evaluation formulas are employed:

1. **Classification Accuracy:** The overall proportion of correctly classified instances across all $K = 8$ classes:

$$\text{Accuracy} = \frac{\sum_{k=1}^K \text{TP}_k}{N_{\text{test}}} = \frac{\text{Correct Predictions}}{\text{Total Predictions}} \quad (34)$$

2. **Precision (Positive Predictive Value):** The fraction of documents assigned to specialty k that actually belong to that specialty:

$$\text{Precision}_k = \frac{\text{TP}_k}{\text{TP}_k + \text{FP}_k} \quad (35)$$

3. **Recall (Sensitivity):** The fraction of documents actually belonging to specialty k that were successfully retrieved by the classifier:

$$\text{Recall}_k = \frac{\text{TP}_k}{\text{TP}_k + \text{FN}_k} \quad (36)$$

4. **F1-Score:** The harmonic mean of Precision and Recall, ensuring that neither precision nor recall is sacrificed:

$$\text{F1-Score}_k = 2 \cdot \frac{\text{Precision}_k \cdot \text{Recall}_k}{\text{Precision}_k + \text{Recall}_k} = \frac{2 \cdot \text{TP}_k}{2 \cdot \text{TP}_k + \text{FP}_k + \text{FN}_k} \quad (37)$$

5. **Macro-Averaged F1-Score (Macro F1):** The unweighted arithmetic mean of F1-scores across all K clinical specialties:

$$\text{Macro F1} = \frac{1}{K} \sum_{k=1}^K \text{F1-Score}_k \quad (38)$$

In medical classification, Macro F1 is the definitive indicator of model health: it assigns equal weight ($1/8 = 12.5\%$) to every clinical department, preventing dominant categories (e.g., Cardiovascular, Orthopedics) from masking poor recall on smaller departments (e.g., Ophthalmology, ENT).

5.3 CONFUSION MATRIX

The confusion matrix $\mathbf{C} \in \mathbb{R}^{K \times K}$ provides an exact visual diagnostic of true versus predicted label assignments, where entry $C_{i,j}$ represents the number of records belonging to true clinical specialty i that were predicted as specialty j . The diagonal elements $C_{i,i}$ denote correct positive identifications (True Positives), while off-diagonal elements indicate clinical misclassifications.

Figure 15 illustrates the 8×8 confusion matrix generated by the Soft Voting Ensemble on the unobserved 333 test records.




Figure 15: Confusion Matrix of the Soft Voting Ensemble across 8 Clinical Specialties on the Held-Out Test Set

The heatmap confirms an intensely pronounced dark diagonal band, indicating high True Positive rates across all eight clinical categories with minimal off-diagonal dispersion.

5.4 RELEVANT METRICS CALCULATED FROM CONFUSION MATRIX

Table 14 presents the granular per-specialty metrics calculated directly from the confusion matrix, including True Positives, False Positives, False Negatives, Precision, Recall, F1-Score, and Support for each medical department.

Clinical Specialty	TP	FP	FN	Precision	Recall	F1-Score	Support
Cardiovascular / Pulmonary	68	9	6	0.88	0.92	0.90	74
ENT - Otolaryngology	16	1	3	0.94	0.84	0.89	19
Gastroenterology	40	6	5	0.87	0.89	0.88	45
Neurology	35	11	10	0.76	0.78	0.77	45
Obstetrics / Gynecology	28	2	3	0.93	0.90	0.92	31
Ophthalmology	17	0	0	1.00	1.00	1.00	17
Orthopedic	63	8	8	0.89	0.89	0.89	71
Urology	28	1	3	0.97	0.90	0.93	31
Overall / Macro Average	295	38	38	0.91	0.89	0.90	333
Weighted Average	—	—	—	0.89	0.89	0.89	333

Table 14: Granular Classification Report and Metrics Derived from the Confusion Matrix

Key analytical observations from Table 14:

- **Ophthalmology:** Achieved a flawless **100% Precision, 100% Recall, and 1.00 F1-score** (17 out of 17 correct). Eye-related vocabulary (e.g., *cataract, lens, cornea, intraocular, phacoemulsification*) forms a distinct, non-overlapping lexical cluster in TF-IDF space.
- **Urology:** Achieved **97% Precision and 0.93 F1-score**, with only 1 false positive across 31 test instances.
- **Obstetrics / Gynecology:** Achieved **93% Precision, 90% Recall, and 0.92 F1-score**, demonstrating high discriminative accuracy for maternal and fetal clinical narratives.
- **Cardiovascular / Pulmonary:** Maintained **92% Recall and 0.90 F1-score** on the largest test cohort (74 records), accurately capturing cardiac catheterization, coronary artery bypass, and pulmonary function narratives.
- **Orthopedic:** Achieved **89% Precision, 89% Recall, and 0.89 F1-score** across 71 test records, successfully resolving musculoskeletal and joint replacement dictations.
- **Neurology:** Scored **76% Precision and 78% Recall**, representing the most challenging classification domain due to terminology overlaps with spinal orthopedic notes.

5.5 PREDICTION STATISTICS OF UNSEEN DATA

An error distribution analysis of the 38 misclassifications on unseen test data reveals clear clinical patterns:

1. **Neurology vs. Orthopedic Boundary Confusion:** Of the 10 false negatives in Neurology, 6 records were misclassified as Orthopedic. Clinical inspection reveals these narratives describe spinal decompression, lumbar disc herniation, and cervical radiculopathy. These procedures involve both neurological structures (spinal cord, nerve roots) and orthopedic biomechanics (vertebrae, bone grafting, spinal fusion hardware), creating a natural overlapping lexical boundary.
2. **Cardiovascular vs. Gastroenterology Boundary:** A small number of cardiovascular patients presenting with epigastric discomfort and gastrointestinal reflux symptoms were mapped to Gastroenterology before cardiac catheterization confirmed coronary artery ischemia.
3. **Prediction Confidence Distribution:** Across all 333 unseen test predictions, the Soft Voting Ensemble demonstrated high confidence:
 - 78.4% of predictions exhibited an ensemble confidence probability exceeding $P \geq 0.85$.
 - 14.1% exhibited confidence between $0.65 \leq P < 0.85$.
 - Only 7.5% exhibited marginal confidence ($P < 0.65$), predominantly clustered around spinal and neuro-orthopedic edge cases.

5.6 DISCUSS YOUR FINDINGS AND CONCLUSIONS

The experimental findings demonstrate that high-performance clinical specialty classification is achievable using classical, computationally efficient machine learning models when domain-specific data engineering principles are rigorously enforced:

1. **Resolution of the Initial Performance Bottleneck:** The initial models evaluated on the raw 40-class dataset suffered from severe failure modes: Random Forest achieved only 7.85% accuracy due to an unconstrained 310,298-feature space, and Logistic Regression

defaulted to 22.96% by predicting only the dominant "Surgery" class. By constraining the TF-IDF feature space to 8,000 informative n-grams and eliminating non-specialty administrative document formats (SOAP notes, Discharge Summaries), accuracy increased dramatically from 27% to **88.59%**.

2. **Efficacy of the Ensemble Paradigm:** Individual algorithms possess unique geometric properties: Linear SVM constructs optimal hyperplanes, Logistic Regression provides smooth log-odds, Random Forest models feature interactions, and Naive Bayes leverages word co-occurrences. The Soft Voting Ensemble combines these diverse paradigms, achieving **87.99% accuracy** and **90.41% macro precision** while significantly reducing prediction variance.
3. **Computational Efficiency and Real-Time Viability:** Unlike deep transformer architectures (such as BioBERT or ClinicalBERT) that require expensive GPUs, extensive memory, and high inference latency, the entire 4-model ensemble trains in **under 3 seconds** on standard CPU hardware and occupies only **11 MB of disk storage**. Inference latency on single patient notes is under 5 milliseconds.
4. **Final Conclusion:** The proposed pipeline successfully classifies unstructured medical transcriptions across 8 major clinical specialties with an overall accuracy of ****88.59%**** and a Macro F1-score of ****0.90****, vastly exceeding the project requirement of 75.00%. The resulting lightweight models are ideally suited for automated patient triage, electronic health record routing, and real-time clinical decision support systems.

References

- [1] A. M. Omar, M. Elhoseny, and A. M. Hassanien, “Clinical text classification with word representation features and machine learning algorithms,” *International Journal of Online and Biomedical Engineering (iJOE)*, vol. 19, no. 4, pp. 40–54, 2023.
- [2] H. Zhang, D. Zhu, H. Tan, M. Shafiq, and Z. Gu, “Medical specialty classification based on semiadversarial data augmentation,” *Hindawi Computational Intelligence and Neuroscience*, vol. 2023, pp. 1–14, 2023.
- [3] A. E. Blanchard, S. Gao, H.-J. Yoon, and G. Tourassi, “A keyword-enhanced approach to handle class imbalance in clinical text classification,” *IEEE Journal of Biomedical and Health Informatics*, vol. 26, no. 6, pp. 2796–2803, Jun. 2022.
- [4] T. Boyle, “Medical transcriptions dataset,” <https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions>, 2019, accessed: 2026-08-03.