

1. INTRODUCTION

1.1 Background

In contemporary healthcare systems, the widespread adoption of Electronic Health Record (EHR) and Electronic Medical Record (EMR) systems has resulted in the generation of massive volumes of clinical documents. Healthcare providers generate extensive unstructured text narratives daily, including admission notes, consultation records, operative reports, discharge summaries, and diagnostic evaluations. Categorizing clinical text into specific medical specialties (e.g., Cardiology, Neurology, Orthopedics, Gastroenterology, Surgery) is a fundamental prerequisite for organizing health databases, routing clinical files, and supporting decision systems.

Traditionally, this categorization relies on manual review by medical coders and health information management staff. As patient volumes increase, manual document indexing becomes increasingly bottlenecked, leading to administrative delays and human classification errors. Natural Language Processing (NLP) and Machine Learning (ML) present viable computational paradigms for automating clinical text categorization, enabling rapid and consistent medical specialty identification.

1.2 Problem Statement

Manual categorization of medical transcription reports into clinical specialties is time-consuming, labor-intensive, and highly prone to inter-annotator variability. Unstructured clinical notes exhibit complex linguistic characteristics, including domain-specific jargon, non-standard medical abbreviations, variable document lengths, and semantic ambiguity. Furthermore, real-world clinical text datasets suffer from severe class imbalance, where dominant medical specialties contain thousands of records, whereas specialized sub-disciplines contain very few. Existing systems lack computationally efficient mechanisms to handle high-dimensional clinical text spaces while remaining lightweight enough for real-time web deployment.

1.3 Need for the System

Automating medical specialty classification yields key operational and clinical benefits:

- *Streamlined Workflow*: Eliminates manual tagging backlogs, allowing clinical reports to

be indexed and routed immediately upon dictation.

- *Enhanced Information Retrieval*: Enables clinicians to query specialty-specific diagnostic records rapidly during emergency care.
- *Reduced Administrative Overhead*: Allows medical records staff to focus on quality assurance rather than manual document sorting.
- *Scalable Health Analytics*: Provides standardized metadata necessary for hospital resource allocation and epidemiological research.
- *Real-Time Accessibility*: Proposes an intuitive web interface for healthcare professionals to upload reports and receive specialty predictions.

1.4 Objectives

The primary objective of this project is to develop and deploy an automated Medical Specialty Classification system using TF-IDF feature representations and an Ensemble Machine Learning architecture. The specific objectives are:

1. To conduct Exploratory Data Analysis (EDA) on the Kaggle Medical Transcriptions (MTSamples) dataset.
2. To build an NLP preprocessing pipeline incorporating lowercasing, punctuation removal, tokenization, stop-word removal, and lemmatization.
3. To extract numerical text features using Term Frequency–Inverse Document Frequency (TF-IDF) vectorization.
4. To train and evaluate candidate machine learning models (Logistic Regression, Support Vector Machine, Random Forest, Multinomial Naive Bayes).
5. To construct a Voting Ensemble Classifier combining predictions from candidate base models.
6. To deploy a real-time, interactive web application using Flask for medical document classification.

1.5 Scope

The scope of this project encompasses:

- *Dataset Scope*: Evaluates 4,999 clinical transcription records across 40 medical specialties from the Kaggle MTSamples dataset.
- *Feature Scope*: Focuses on TF-IDF vectorization derived from the primary narrative column (transcription).
- *Algorithmic Scope*: Evaluates classical ML models (LR, SVM, RF, MNB) and a Voting Ensemble Classifier optimized for CPU deployment.
- *Deployment Scope*: Proposes a lightweight Flask web interface for single-document and batch specialty prediction.

1.6 Project Overview

The project follows a structured engineering workflow:

1. *Data Ingestion & EDA*: Ingesting raw MTSamples CSV data and analyzing feature distributions.
2. *Text Preprocessing*: Removing noise, tokenizing, filtering stop-words, and lemmatizing text.
3. *Feature Engineering*: Transforming clean text into sparse numerical vectors via TF-IDF.
4. *Model Development*: Training baseline classifiers and hyperparameter tuning.
5. *Ensemble Integration*: Aggregating base estimator probabilities via soft/hard voting.
6. *Web Deployment*: Serializing model artifacts with Joblib and serving via a proposed Flask web application.

1.7 Expected Outcome

The expected deliverables include:

- A clean, preprocessed clinical text dataset optimized for NLP tasks.

- A proposed Voting Ensemble Classifier targeting robust performance across 40 medical specialties.
- A planned lightweight Flask web application to provide real-time specialty classification.
- Complete technical documentation adhering to MCA departmental standards.

2. SUPPORTING LITERATURE

2.1 Literature Review

A systematic review of contemporary literature was conducted to evaluate existing methodological approaches, feature representations, classification algorithms, and dataset challenges in medical text categorization. The review relies strictly on approved reference literature in clinical natural language processing.

Paper 1 – Omar AM, Elhoseny M, Hassanien AM. Clinical text classification with word representation features and machine learning algorithms. International Journal of Online and Biomedical Engineering (iJOE). 2023.

The paper titled “Clinical Text Classification with Word Representation Features and Machine Learning Algorithms” evaluates the impact of different feature extraction techniques and machine learning algorithms on Electronic Medical Record (EMR) narratives. The authors introduced a four-phase classification framework consisting of text preprocessing, word vectorization, feature reduction, and supervised classification.

To model clinical text narratives, the study compared classical bag-of-words and TF-IDF representations against dense neural embeddings (Word2Vec) across multiple classifiers, establishing baseline benchmarks for clinical text processing. The evaluated classifiers include Logistic Regression, Support Vector Machine (SVM), Naive Bayes, and k-Nearest Neighbors (k-NN).

The empirical results demonstrated that combining Word2Vec embeddings with a k-NN classifier achieved the highest accuracy of 92%. Dense semantic embeddings consistently outperformed sparse unigram Bag-of-Words representations across clinical text categories. However, distance-based k-NN introduces significant computational latency during inference, rendering it inefficient for massive production databases. Furthermore, the study relied on single classifiers without exploring ensemble learning strategies or real-time web deployment frameworks. Crucially, Principal Component Analysis (PCA) was evaluated exclusively in Paper 1 for feature reduction and is not included in our proposed system architecture.

TITLE	Clinical Text Classification with Word Representation Features and Machine Learning Algorithms [1]
AREA OF WORK	Clinical Natural Language Processing / EMR Text Vectorization
DATASET	Electronic Medical Record (EMR) Clinical Transcriptions Dataset
METHODOLOGY / STRATEGY	4-Phase Framework: (1) Preprocessing, (2) Vectorization (BOW, TF-IDF, Word2Vec), (3) Feature Reduction (PCA), (4) Classification
ALGORITHM	Logistic Regression, Support Vector Machine (SVM), Naive Bayes, k-NN, PCA
RESULT / PRECISION	Word2Vec + k-NN achieved the highest classification accuracy of 92%
ADVANTAGES	Proves dense semantic embeddings capture medical context better than unigram BOW
LIMITATIONS	Distance-based k-NN scales poorly; lacks ensemble learning and web UI; PCA used only in literature
RELEVANCE TO MY PROJECT	Establishes the standard 4-phase text preprocessing and feature extraction pipeline
FUTURE PROPOSAL	Develop lightweight ensemble classifiers suitable for real-time web deployment

Table 1: Summary of Paper 1

Paper 2 – Zhang H, Zhu D, Tan H, Shafiq M, Gu Z. Medical specialty classification based on semiadversarial data augmentation. Hindawi Computational Intelligence and Neuroscience. 2023.

The research article titled “Medical Specialty Classification Based on Semiadversarial Data Augmentation” addressed data scarcity and category imbalance in medical specialty classification. The authors developed Semiadversarial Data Augmentation (SemiADA), generating synthetic samples along adversarial attack trajectories to populate sparse decision regions.

Additionally, a Probabilistic Information (PI) layer was introduced post-Softmax to recalculate confidence based on domain noun importance. The paper evaluated a filtered subset of the Kaggle Medical Specialty Classification dataset restricted to 18 medical specialty classes. In contrast, our project utilizes the complete Kaggle Medical Transcriptions (MTSamples) dataset spanning 40 medical specialties.

The proposed SemiADA framework achieved an average performance improvement of 14.9% in accuracy and F1 score over baseline benchmarks. The study proved that domain-specific nouns carry the highest discriminative signal for medical specialty assignment. However, generating synthetic text along adversarial attack paths incurs extreme computational overhead. Fine-tuning heavy transformer models like BERT requires dedicated GPU infrastructure, making real-time inference difficult in resource-constrained web environments.

TITLE	Medical Specialty Classification Based on Semiadversarial Data Augmentation [2]
AREA OF WORK	Clinical Text Data Augmentation / Class Imbalance Handling
DATASET	Kaggle Medical Specialty Classification Dataset (Subset restricted to 18 classes)
METHODOLOGY / STRATEGY	Semiadversarial sample generation along attack paths + Probabilistic Information (PI) noun-weighted Softmax recalculation
ALGORITHM	SemiADA, BERT (Transformer), Softmax with PI Noun Layer
RESULT / PRECISION	Achieved an average 14.9% boost in Accuracy and F1 score over baseline benchmarks
ADVANTAGES	Effectively handles extreme class imbalance using geometry-aware synthetic samples
LIMITATIONS	Extreme computational latency; fine-tuning BERT demands high-end GPU hardware
RELEVANCE TO MY PROJECT	Identifies class imbalance in MTSamples data; motivates TF-IDF + Ensemble ML alternative
FUTURE PROPOSAL	Combine TF-IDF term weighting with ensemble learning to classify 40 specialties on standard CPUs

Table 2: Summary of Paper 2

Paper 3 – Blanchard AE, Gao S, Yoon HJ, Tourassi G. A keyword-enhanced approach to handle class imbalance in clinical text classification. IEEE Journal of Biomedical and Health Informatics. 2022.

Blanchard et al., writing in the *IEEE Journal of Biomedical and Health Informatics*, focused on preventing deep neural models from overfitting on rare clinical categories in long pathology narratives. The authors introduced a data-centric strategy by injecting short class-associated

keyword sequences into mini-batches during neural network training, formulating a combined loss function

$$L = L_{\text{docs}} + \alpha L_{\text{key}}$$

The study utilized the National Cancer Institute (NCI) Surveillance, Epidemiology, and End Results (SEER) cancer pathology registry dataset. The approach evaluated Convolutional Neural Networks (CNNs), keyword-constrained loss functions, and standard data resampling baselines (oversampling and undersampling).

Injecting keyword constraints significantly improved macro F1 score and recall on rare classes. The study demonstrated that overall micro-accuracy is a deceptive metric in clinical NLP because high overall scores can mask near-zero recall on rare clinical conditions. However, the approach requires manual curation of comprehensive keyword dictionaries by clinical domain experts for every target category, limiting automated scalability to new medical domains.

TITLE	A Keyword-Enhanced Approach to Handle Class Imbalance in Clinical Text Classification [3]
AREA OF WORK	Imbalanced Clinical Narrative Classification / Cancer Registry NLP
DATASET	NCI SEER Surveillance Cancer Pathology Registry Reports Dataset
METHODOLOGY / STRATEGY	Ingesting keyword sequences into mini-batches during training to constrain the neural loss function
ALGORITHM	Convolutional Neural Networks (CNNs), Keyword-Constrained Loss Optimization
RESULT / PRECISION	Significantly improved macro F1 score and recall on rare classes without degrading majority performance
ADVANTAGES	Prevents deep models from memorizing noisy tokens in long text documents with rare labels
LIMITATIONS	Requires manual domain expert curation of keyword dictionaries for every target category
RELEVANCE TO MY PROJECT	Demonstrates that keyword/n-gram features preserve rare class signals, supporting TF-IDF
FUTURE PROPOSAL	Utilize automated TF-IDF n-gram feature extraction without requiring manual dictionary creation

Table 3: Summary of Paper 3

2.2 Literature Review Summary

Table 4 presents a concise summary of the reviewed literature alongside our proposed system.

Paper	Area	Dataset	Algorithm	Advantages	Relation to Proposed Work
Omar et al. [1]	Clinical Vectorization	EMR Text	LR, SVM, NB, k-NN	High accuracy (92%)	Establishes 4-phase preprocessing pipeline
Zhang et al. [2]	Data Augmentation	MTSamples (18 classes)	SemiADA, BERT	+14.9% accuracy gain	Highlights MTSamples class imbalance
Blanchard et al. [3]	Class Imbalance	NCI SEER Pathology	CNN + Keyword Loss	High macro F1 on rare classes	Justifies TF-IDF term weighting
Proposed Work	Specialty Categorization	MTSamples (40 classes)	TF-IDF + Voting Ensemble	Fast, lightweight, web UI	Complete MCA Project System

Table 4: Overall Summary of all paper

Evolution of Research

Clinical text processing research has transitioned from basic Bag-of-Words word counting toward neural word embeddings (Word2Vec), transformer architectures (BERT), and complex adversarial data augmentation (SemiADA). While deep neural networks achieve strong performance, their high computational cost and memory footprints prevent real-time deployment in standard healthcare web environments.

Identified Research Gaps

The literature review reveals three primary research gaps:

1. **High Computational Overhead:** Existing deep learning systems require dedicated GPUs and introduce substantial inference latency.
2. **Lack of Real-Time Web Deployment:** Prior studies focus exclusively on offline evalua-

tion metrics without building accessible user interfaces.

3. **Single Model Reliance:** Most prior works rely on single classifiers rather than leveraging multi-model ensemble diversity.

Motivation for Proposed System

Our project addresses these research gaps by combining TF-IDF vectorization with a Voting Ensemble Classifier (aggregating Logistic Regression, Linear SVM, Random Forest, and Multinomial Naive Bayes) intended for deployment via a lightweight Flask web application. This architecture is designed to target accurate classification across all 40 medical specialties in the full MTSamples dataset while aiming for efficient CPU inference for real-world clinical use.

2.3 Findings and Proposals

Why TF-IDF?

Term Frequency–Inverse Document Frequency (TF-IDF) is selected because it evaluates word importance relative to the entire corpus. In medical transcriptions, common clinical terms appear across all documents, whereas domain-specific keywords (e.g., "endoscopy", "myocardial", "arthroscopy") appear frequently within specific specialties. TF-IDF automatically scales down uninformative common words while assigning high numerical weights to discriminative medical terms.

Why Ensemble Machine Learning?

Ensemble machine learning combines predictions from diverse base models (linear, tree-based, margin-based, and probabilistic). Base models exhibit different decision boundaries and error profiles. A Voting Ensemble averages these individual prediction probabilities, effectively reducing prediction variance, overcoming single-model bias, and producing robust specialty predictions.

Why Flask?

Flask is a lightweight, WSGI-compliant Python web framework. It provides minimal overhead, seamless integration with Python machine learning libraries (Scikit-learn, Joblib), and fast route execution, making it ideal for hosting real-time inference endpoints.

Expected Improvements

The proposed system is expected to deliver fast document classification, stable predictive performance across imbalanced medical categories, and an intuitive web interface for clinical document indexing.

3. SYSTEM ANALYSIS

The Kaggle Medical Transcriptions (MTSamples) dataset, sourced from Kaggle [4], provides clinical narrative dictations across multiple healthcare sub-disciplines, enabling comprehensive evaluation of natural language processing and machine learning models for medical document categorization.

Dataset Link: <https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions>

Size: The dataset comprises 4,999 clinical records spanning 40 distinct medical specialties.

Data Structure: It contains six attributes including row index (Unnamed: 0), short report abstract (description), target medical department (medical_specialty), sample dictation title (sample_name), main clinical narrative (transcription), and extracted key terms (keywords).

Sparsity: The primary text column (transcription) is highly complete, containing 33 missing records (0.66%), which are dropped during preprocessing. The target label (medical_specialty) contains zero missing values.

3.1 Analysis of dataset

3.1.1 About the dataset

a) Dataset Overview

The Kaggle Medical Transcriptions dataset provides authentic clinical narratives dictating patient histories, diagnostic evaluations, physical examinations, and surgical procedures:

- Total Records: 4,999 clinical reports.
- Total Features: 5 independent variables + 1 target variable (medical_specialty).
- Data Types: Textual narratives (transcription, description, keywords), categorical labels (medical_specialty), and integer index.
- Missing Values: Present in description (6 records), transcription (33 records), and keywords (1,068 records).

Table 5 summarizes the core dataset attributes.

Attribute	Details
Dataset Name	Kaggle Medical Transcriptions (MTSamples) Dataset
Source URL	https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions
Number of Records	4,999 clinical reports
Number of Columns	6 fields (Unnamed: 0, description, medical_specialty, sample_name, transcription, keywords)
Input Variable	transcription (Primary clinical narrative text)
Target Variable	medical_specialty (40 unique clinical specialties)
Advantages	Authentic clinical dictations; complete target label annotations (0 missing labels); broad multi-class domain coverage
Limitations	Severe class imbalance across categories; 33 missing transcriptions; 1,068 missing keyword entries; variable text length

Table 5: Kaggle Medical Transcriptions (MTSamples) Dataset Summary

3.1.2 Explore the dataset

A statistical audit was conducted on the dataset fields. Table 6 describes each column, its data type, and missing value counts.

Column Name	Data Type	Description	Missing Count
Unnamed: 0	Integer	Record Identifier / Row Index	0
description	Text	Short summary of clinical report	6
medical_specialty	Categorical	Target specialty label (40 classes)	0
sample_name	Text	Report title / dictation sample title	0
transcription	Text	Primary clinical narrative (Input)	33
keywords	Text	Extracted key terms	1,068

Table 6: Dataset Column Descriptions and Missing Value Audit

Table 7 lists the record counts for the top 10 majority specialties alongside the 10 rarest minority specialties.

Top 10 Specialties	Count	Bottom 10 Specialties	Count
Surgery	1,103	Cosmetic / Plastic Surgery	27
Consult - History and Phy.	516	Pediatric / Neonatal	22
Cardiovascular / Pulmonary	372	Sleep Medicine	20
Orthopedic	355	Bariatrics	18
Radiology	273	Speech - Language	9
General Medicine	259	Lab Medicine - Pathology	8
Gastroenterology	230	Allergy / Immunology	7
Neurology	223	Hospice - Palliative Care	6
SOAP / Chart / Progress Notes	166	Executive Evaluation	2
Obstetrics / Gynecology	160	Autopsy	2

Table 7: Distribution of Top 10 Majority and Bottom 10 Minority Medical Specialties

3.2 Data Preprocessing

3.2.1 Data Cleaning

Text preprocessing is essential in clinical NLP to remove noisy elements, normalize medical vocabulary, and transform raw narrative text into clean input suitable for mathematical vectorization. The cleaning pipeline executes eight sequential operations:

1. *Removing Missing Values*: Records containing empty or null transcription fields (33 records) are dropped from the dataset.
2. *Removing Duplicates*: Verification ensures no duplicate clinical reports exist in the corpus.
3. *Lowercasing*: All narrative text is converted to lowercase to standardize character representations.
4. *Removing Punctuation*: Commas, periods, semicolons, and special punctuation symbols are stripped.
5. *Removing Special Characters*: Numeric digits and non-alphanumeric symbols are removed.
6. *Tokenization*: Free-text narratives are segmented into individual word tokens.
7. *Stop-Word Removal*: Generic English stop-words are filtered out while preserving clinical vocabulary.
8. *Lemmatization*: Word tokens are mapped to their canonical base form using the WordNet Lemmatizer.

Table 8 provides concrete text transformation examples for each step.

Preprocessing Step	Input Text Example	Processed Output Example
Raw Input Text	"SUBJECTIVE: Patient suffers from acute chest pain!"	"SUBJECTIVE: Patient suffers from acute chest pain!"
1. Missing Values	Checked: Non-null text	Retained
2. Duplicates	Checked: Unique entry	Retained
3. Lowercasing	"SUBJECTIVE: Patient suffers..."	"subjective: patient suffers from acute chest pain!"
4. Punctuation	"subjective: patient suffers..."	"subjective patient suffers from acute chest pain"
5. Special Chars	"chest pain! 100%"	"chest pain"
6. Tokenization	"patient suffers pain"	['patient', 'suffers', 'pain']
7. Stop-Words	['patient', 'suffers', 'from', 'pain']	['patient', 'suffers', 'pain']
8. Lemmatization	['patient', 'suffers', 'pain']	['patient', 'suffer', 'pain']

Table 8: Step-by-Step Clinical Text Preprocessing Examples

3.2.2 Analysis of Feature Variables

The transcription field is selected as the sole feature variable because it contains full clinical narratives describing patient symptoms, diagnostic findings, and surgical procedures. Meta-data fields (description, keywords) are excluded due to missingness. Numerical features are extracted using TF-IDF vectorization, which scales term frequency by inverse document frequency to emphasize discriminative medical terms.

3.2.3 Analysis of Class Variables

The target variable `medical_specialty` contains 40 categories. Class imbalance (Surgery = 1,103 vs Autopsy = 2) is addressed using class weighting in estimators and voting ensemble aggregation to maintain balanced predictive performance.

3.3 Data Visualization

Exploratory Data Analysis (EDA) was performed on the Kaggle Medical Transcriptions (MT-Samples) dataset to analyze missing attribute distributions, class target distributions, document length characteristics, and clinical vocabulary frequency.

Missing Value Analysis

Figure 1 illustrates the distribution of missing values across dataset attributes. As shown, the primary target attribute (`medical_specialty`) contains zero missing values. The primary input text narrative (`transcription`) contains 33 missing records, while the auxiliary keywords attribute exhibits 1,068 missing records.




Figure 1: Missing Values by Dataset Attribute

Top 10 Medical Specialties

Figure 2 displays the record counts for the top 10 medical specialties in the dataset. Surgery represents the single largest category with 1,103 clinical reports, followed by Consult - History and Phy. (516 reports), Cardiovascular / Pulmonary (372 reports), and Orthopedic (355 reports).




Figure 2: Top 10 Medical Specialties by Number of Transcriptions

Complete Specialty Distribution

Figure 3 illustrates the full class distribution across all 40 medical specialties in the MTSamples corpus. The chart demonstrates severe class imbalance, ranging from majority classes such as Surgery (1,103 records) down to highly specialized minority classes such as Hospice - Palliative Care, Allergy / Immunology, Lab Medicine - Pathology, and Autopsy with under 10 records each.




Figure 3: Distribution of Medical Specialties

Transcription Length Distribution

Figure 4 displays the word count distribution for clinical narratives, fitted with a Kernel Density Estimate (KDE) curve. The majority of medical transcriptions range between 200 and 600 words, exhibiting a right-skewed distribution with extended narratives reaching up to 3,000 words.




Figure 4: Distribution of Clinical Transcription Length

3.4 Analysis of Algorithms

3.4.1 Detailed Study Report of the Proposed Algorithms

TF-IDF Vectorizer

Term Frequency–Inverse Document Frequency (TF-IDF) is a statistical vectorization method used to convert unstructured clinical text into high-dimensional numerical feature vectors. It evaluates the importance of a word relative to an individual report as well as the entire corpus.

Process Flow: TF-IDF computes the mathematical product of local term frequency $TF(t, d)$ and logarithmic inverse document frequency

$$IDF(t, D) = \log(N/DF(t))$$

Words appearing frequently across all clinical reports receive lower weights, while specialized terms receive higher numerical weights.

Advantages: Computationally efficient, scales to large vocabularies, and highlights discriminative medical keywords without requiring GPU infrastructure.

Limitations: Treats text as an unordered bag-of-words and ignores sequential context or semantic word relationships.

Reason for Selection: Provides high-dimensional sparse numerical inputs optimized for linear classifiers.

Role in Proposed System: Serves as the primary feature extraction engine, converting lemmatized clinical narratives into sparse numerical feature matrices for classifier training.

Logistic Regression

Logistic Regression is a linear classification algorithm extended to multi-class clinical text categorization through multinomial Softmax regression. It provides a strong, interpretable probabilistic baseline for text processing.

Process Flow: Computes linear combinations of input TF-IDF features and applies the Softmax function to convert raw model logits into normalized probability distributions across all 40

medical specialties.

Advantages: Fast training speed, calibrated probability outputs, and strong resistance to overfitting when paired with L2 regularization.

Limitations: Assumes linear decision boundaries in feature space, which may struggle with complex non-linear word interactions.

Reason for Selection: Provides a reliable linear probabilistic baseline for sparse text features.

Role in Proposed System: Operates as a primary linear probabilistic base estimator within the Voting Ensemble Classifier.

Support Vector Machine (SVM)

Support Vector Machines (SVM) are supervised learning classifiers that construct maximum-margin hyperplanes separating target classes in high-dimensional feature space.

Process Flow: Solves convex quadratic optimization to maximize the margin between closest training samples (support vectors) belonging to different medical specialties using linear hinge loss.

Advantages: Highly effective in sparse, high-dimensional TF-IDF feature spaces and resistant to overfitting.

Limitations: Does not generate native probability estimates without Platt scaling and requires higher memory during training.

Reason for Selection: Proven state-of-the-art classifier for high-dimensional TF-IDF feature spaces.

Role in Proposed System: Serves as the primary margin-based linear estimator in the Voting Ensemble, maximizing separation between medical categories.

Random Forest

Random Forest is an ensemble learning method that constructs a multitude of decision trees during training and outputs the majority class vote across individual trees.

Process Flow: Combines bootstrap aggregation (bagging) with random feature subspace se-

lection at each tree node split, measuring impurity reduction via the Gini criterion.

Advantages: Captures complex non-linear term interactions and demonstrates high robustness against noise and outliers.

Limitations: Computationally slower during training on large, sparse high-dimensional matrices.

Reason for Selection: Acts as the non-linear tree-based estimator in the Voting Ensemble, providing architectural model diversity.

Role in Proposed System: Operates as the tree-based estimator in the Voting Ensemble.

Multinomial Naive Bayes

Multinomial Naive Bayes is a probabilistic classifier based on Bayes' Theorem, specifically designed for text classification with word frequency counts.

Process Flow: Computes log-posterior class probabilities

$$P(y|x) \propto P(y) \prod P(x_i|y)$$

using word frequency counts with Laplace smoothing to prevent zero-probability errors.

Advantages: Extremely fast training and inference speeds with low computational overhead.

Limitations: Relies on the strong assumption that feature words are conditionally independent given the class.

Reason for Selection: Industry standard baseline classifier tailored for word count and frequency features.

Role in Proposed System: Operates as a fast probabilistic base estimator within the Voting Ensemble Classifier.

Voting Ensemble Classifier

The Voting Ensemble Classifier is a meta-estimator that aggregates predictions from individual base models (Logistic Regression, SVM, Random Forest, Multinomial Naive Bayes) to yield a

unified prediction.

Process Flow: Computes the average predicted probability across all base estimators

$$P_{\text{ensemble}}(y = c) = \frac{1}{M} \sum P_m(y = c)$$

and assigns the class label corresponding to the maximum aggregated probability.

Advantages: Reduces individual model variance and bias, producing stable predictions across imbalanced clinical specialties.

Limitations: Total inference latency equals the cumulative prediction time of all underlying base models.

Reason for Selection: Designed to improve classification performance by aggregating linear, margin, tree, and probabilistic models.

Role in Proposed System: Serves as the decision engine to be integrated into the Flask web application for real-time specialty categorization.

Activity Diagram

Figure 7 displays the TikZ activity diagram representing the execution flow of the proposed machine learning algorithm.




Figure 7: Activity Diagram of the Proposed Machine Learning Algorithm

Step-by-Step Explanation

The system executes in eight sequential stages:

1. *Dataset Ingestion*: Ingesting raw MTSamples CSV data containing 4,999 medical transcriptions.
2. *Cleaning*: Dropping empty records, removing duplicate entries, and lowercasing narrative text.
3. *Preprocessing*: Removing punctuation, digits, and stop-words, followed by WordNet lemmatization.
4. *TF-IDF Vectorization*: Building sparse numerical feature matrices with sublinear term frequency scaling.
5. *Train-Test Partitioning*: Splitting data into 80% training and 20% testing stratified sets.
6. *Multi-Model Training*: Fitting Logistic Regression, Linear SVM, Random Forest, and Multinomial Naive Bayes.
7. *Ensemble Aggregation*: Combining base estimator outputs into a soft/hard Voting Classifier.
8. *Real-Time Prediction*: Returning top predicted medical specialty with confidence scores.

3.5 Project pipeline

3.5.1 Project pipeline diagram with explanation

Figure 8 presents the engineering project pipeline diagram following a top-down flow.




Figure 8: Engineering Project Pipeline Flow Diagram

Data Phase

The procedure starts with gathering raw clinical data from the Kaggle MTSamples dataset. Preprocessing actions such as addressing missing values, removing noise, and lemmatizing word tokens are implemented. After cleaning, EDA is performed to visualize distributions, identify category counts, and discover patterns for feature extraction.

Model Building Phase

After data preparation, the processed TF-IDF feature matrix will be used to train candidate machine learning models (Logistic Regression, Support Vector Machine, Random Forest, Multinomial Naive Bayes). Predictions from candidate base models will be aggregated into a Voting Ensemble Classifier, and standard classification metrics—precision, recall, and F1 score—will be evaluated to assess model performance across the 40 medical specialties.

Deployment Phase

In the proposed system, the trained Voting Ensemble model will be serialized using Joblib and integrated into a Flask web application. Healthcare professionals will be able to paste or upload clinical notes through the user interface and receive real-time medical specialty predictions.

3.6 Feasibility analysis

3.6.1 Technical Feasibility

The system utilizes standard open-source Python frameworks (Scikit-learn, NLTK, Pandas, Flask) running on standard CPU hardware. No expensive GPU infrastructure is required, establishing technical feasibility.

3.6.2 Economic Feasibility

All development tools, libraries, and datasets are open-source with zero licensing fees. Deploying on existing computer hardware involves no financial expense, establishing economic feasibility.

3.6.3 Operational Feasibility

The proposed web interface is designed to allow non-technical healthcare personnel to upload medical reports and view specialty predictions easily, requiring minimal training and ensuring operational feasibility.

3.6.4 Legal Feasibility

The Kaggle MTSamples dataset consists of public records that do not contain Protected Health Information (PHI). The project utilizes this anonymized public dataset and considers relevant data-privacy requirements, establishing legal feasibility.

3.6.5 Schedule Feasibility

Project milestones are structured systematically across the academic semester, guaranteeing completion within departmental deadlines.

References

- [1] A. M. Omar, M. Elhoseny, and A. M. Hassanien, "Clinical text classification with word representation features and machine learning algorithms," *International Journal of Online and Biomedical Engineering (iJOE)*, vol. 19, no. 4, pp. 40–54, 2023.
- [2] H. Zhang, D. Zhu, H. Tan, M. Shafiq, and Z. Gu, "Medical specialty classification based on semiadversarial data augmentation," *Hindawi Computational Intelligence and Neuroscience*, vol. 2023, pp. 1–14, 2023.
- [3] A. E. Blanchard, S. Gao, H.-J. Yoon, and G. Tourassi, "A keyword-enhanced approach to handle class imbalance in clinical text classification," *IEEE Journal of Biomedical and Health Informatics*, vol. 26, no. 6, pp. 2796–2803, Jun. 2022.
- [4] T. Boyle, "Medical transcriptions dataset," <https://www.kaggle.com/datasets/tboyle10/medicaltranscriptions>, 2019, accessed: 2026-08-03.